

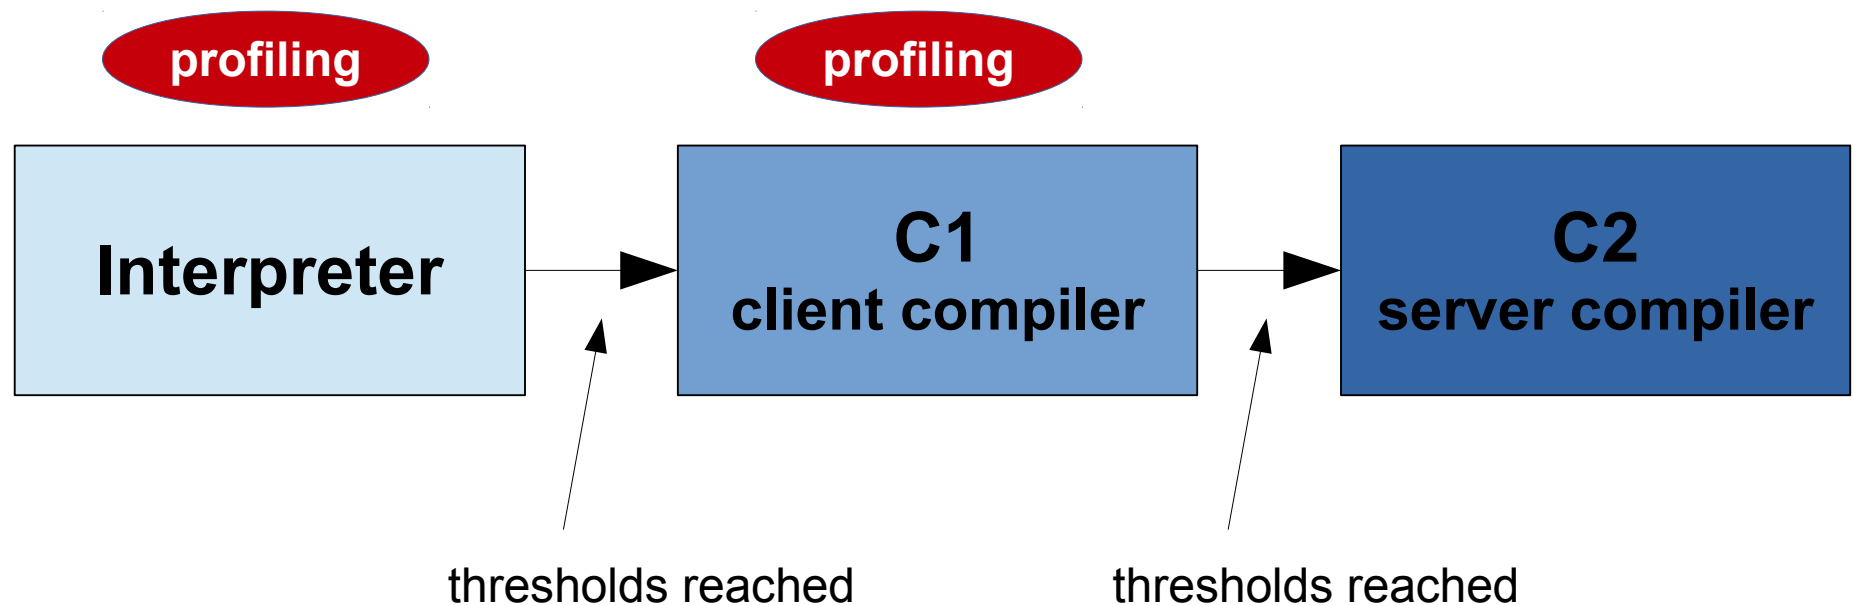
Profile Caching for the Java Virtual Machine

Marcel Mohler, ETH Zurich
Bachelor Thesis

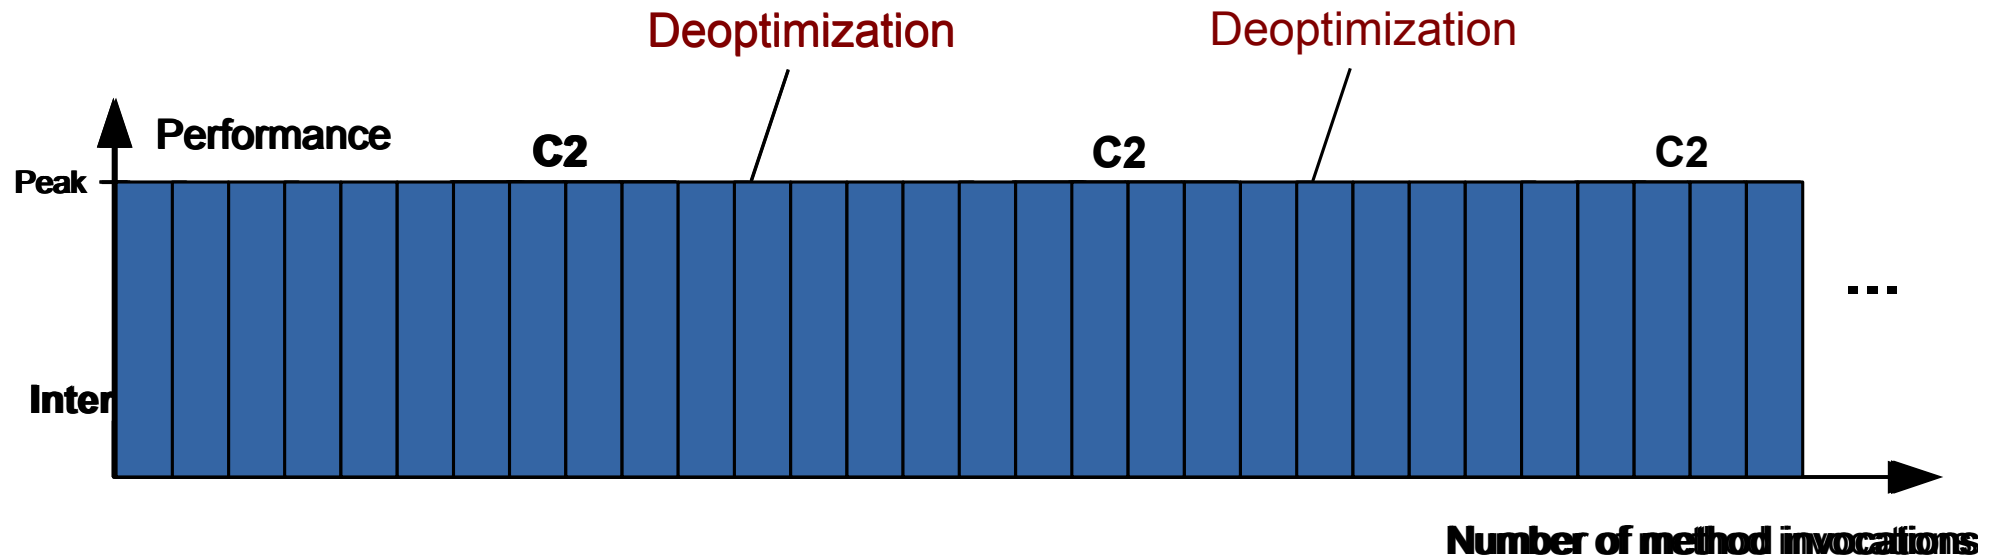
*Supervisors: Zoltan Majo, Oracle
Tobias Hartmann, Oracle*

Prof. Thomas Gross, Laboratory for Software Technology, ETH

Tiered Compilation in HotSpot JVM



Problem: performance fluctuations



Goal: Decrease performance fluctuations

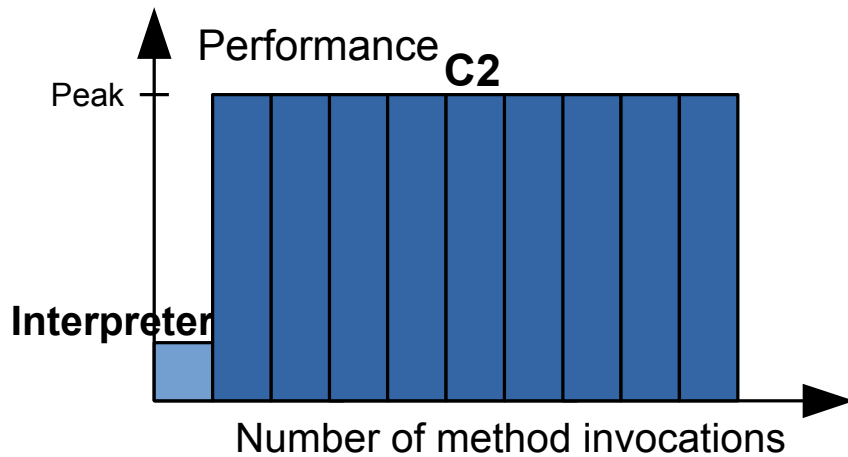
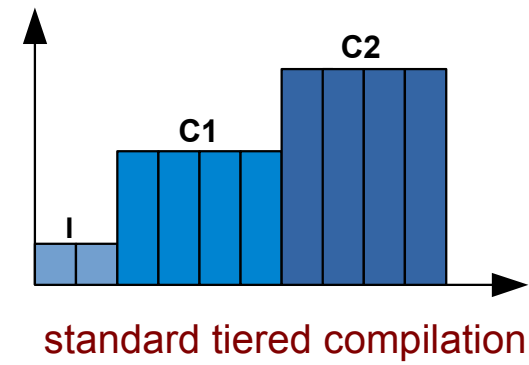
Observation

- Profiles need to be gathered each time the JVM starts
 - Most frequently used methods often do not change
- Idea: cache and reuse the profiles!


Outline

1. Design
2. Implementation
3. Performance evaluation
4. Conclusion

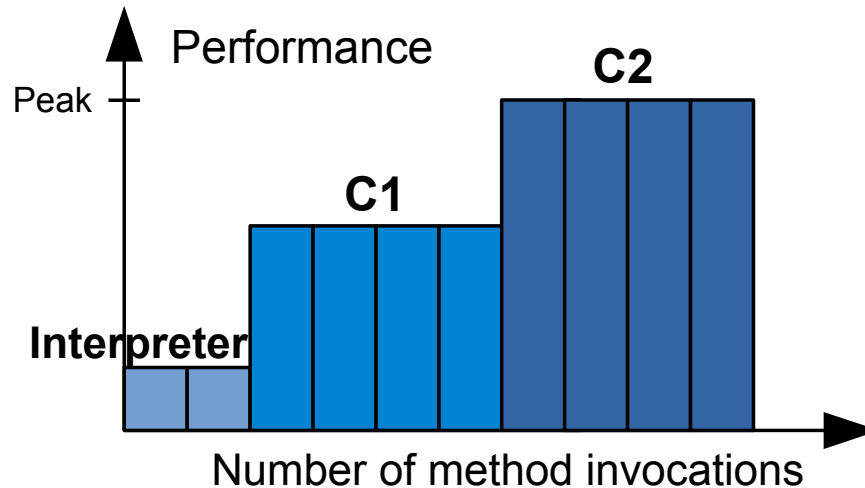
Design: 3 modes



Mode 0

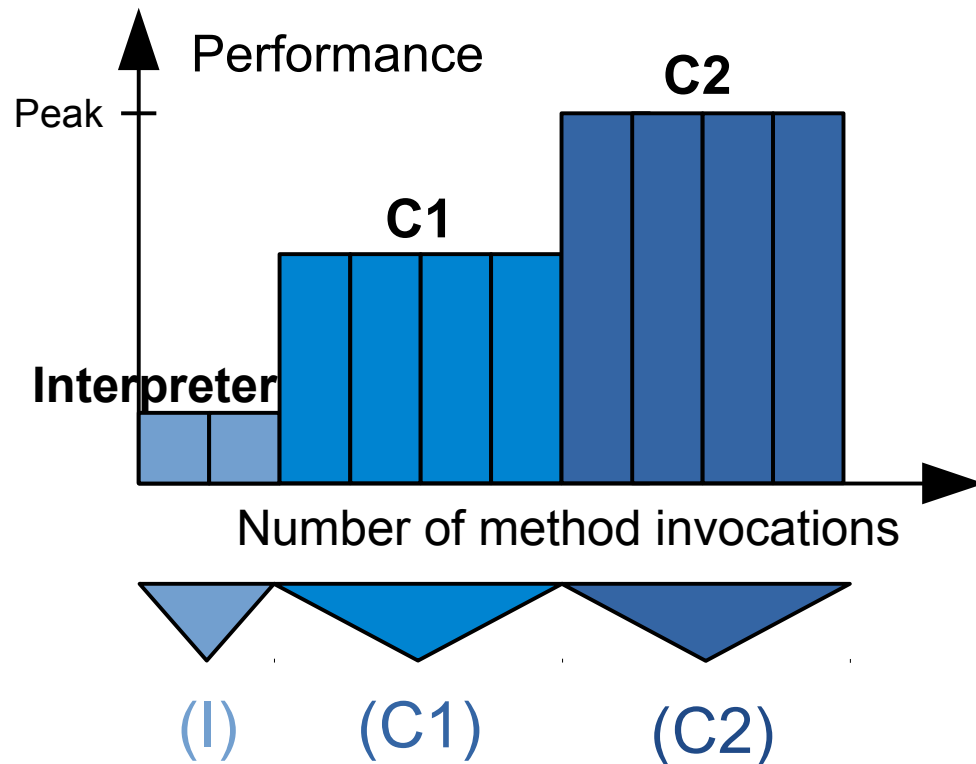


Mode 1



Mode 2

Design: Mode 2



(I) unmodified

(C1) remove profiling code
→ ~30% speedup

(C2) use cached profile
→ better code quality

Implementation

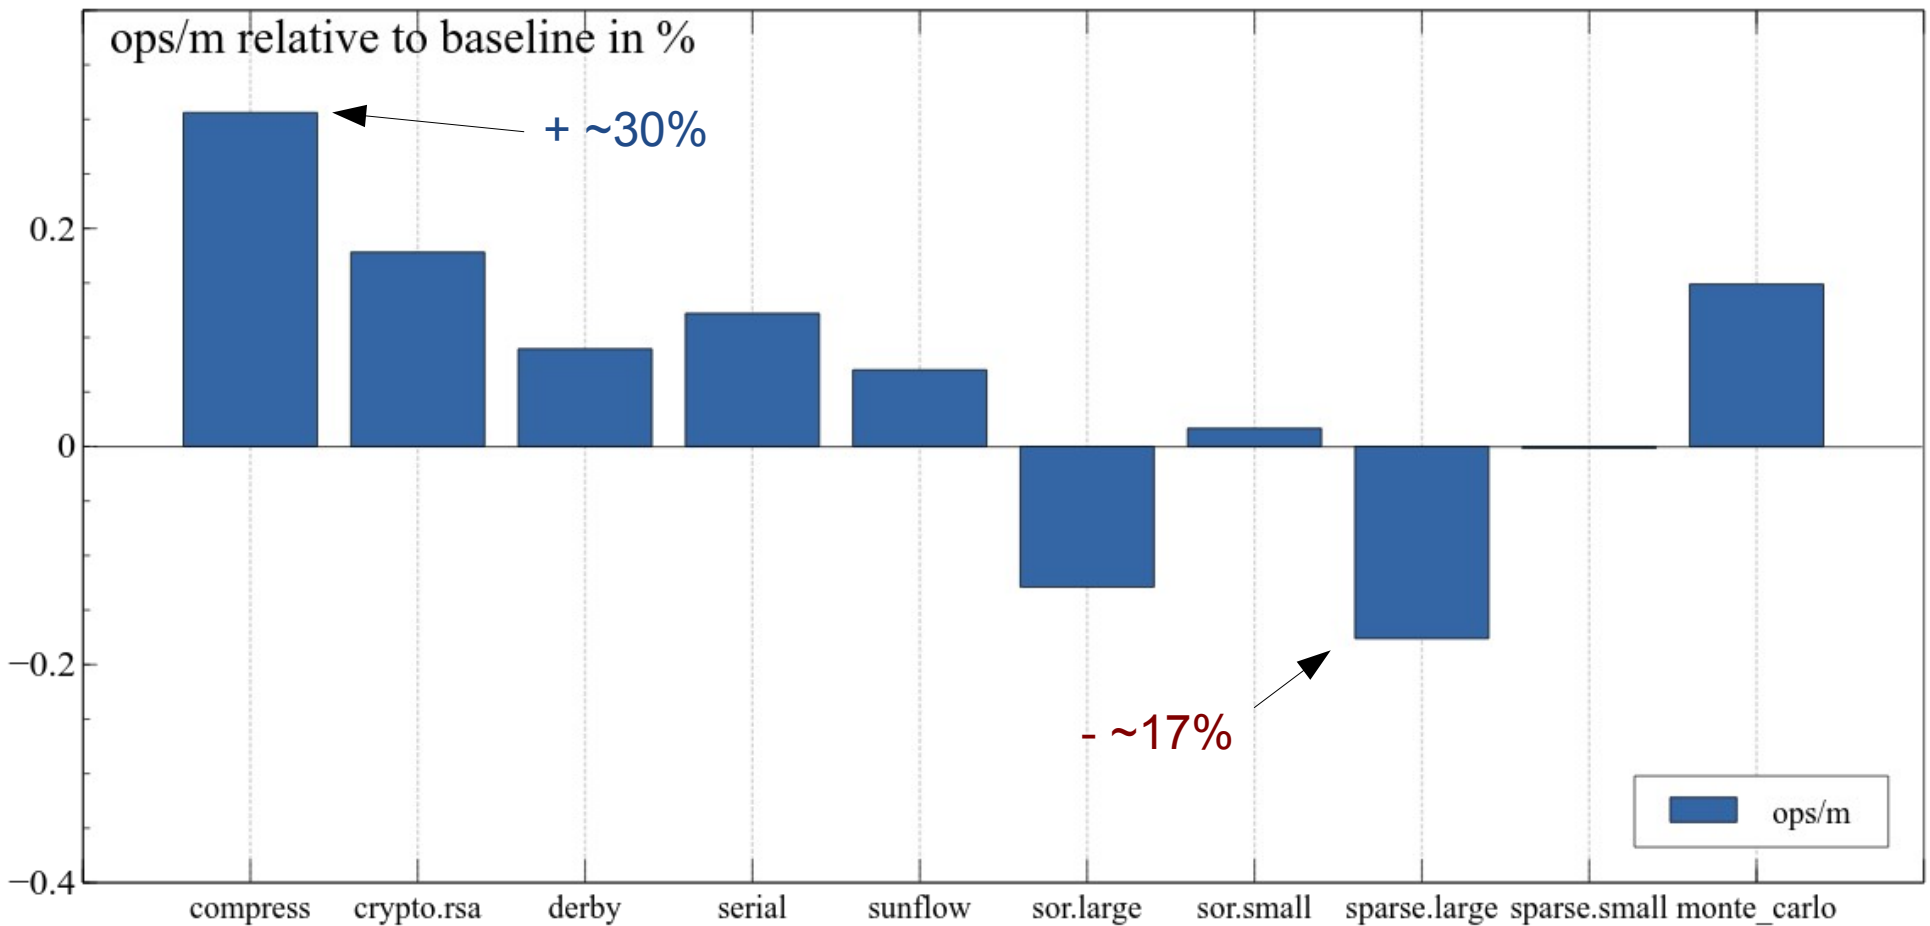
- Based on **existing** compilation replay functionality
- **Configurable** by JVM flags
- Select **all** or an arbitrary **set of** methods

Evaluation

- ETH Data Center Observatory
- Focus on **warmup**
- 2 **benchmark suites**
 - SPECjvm 2008
17/21 individual benchmarks used
 - Google Octane (using Nashorn)
16/17 individual benchmarks used

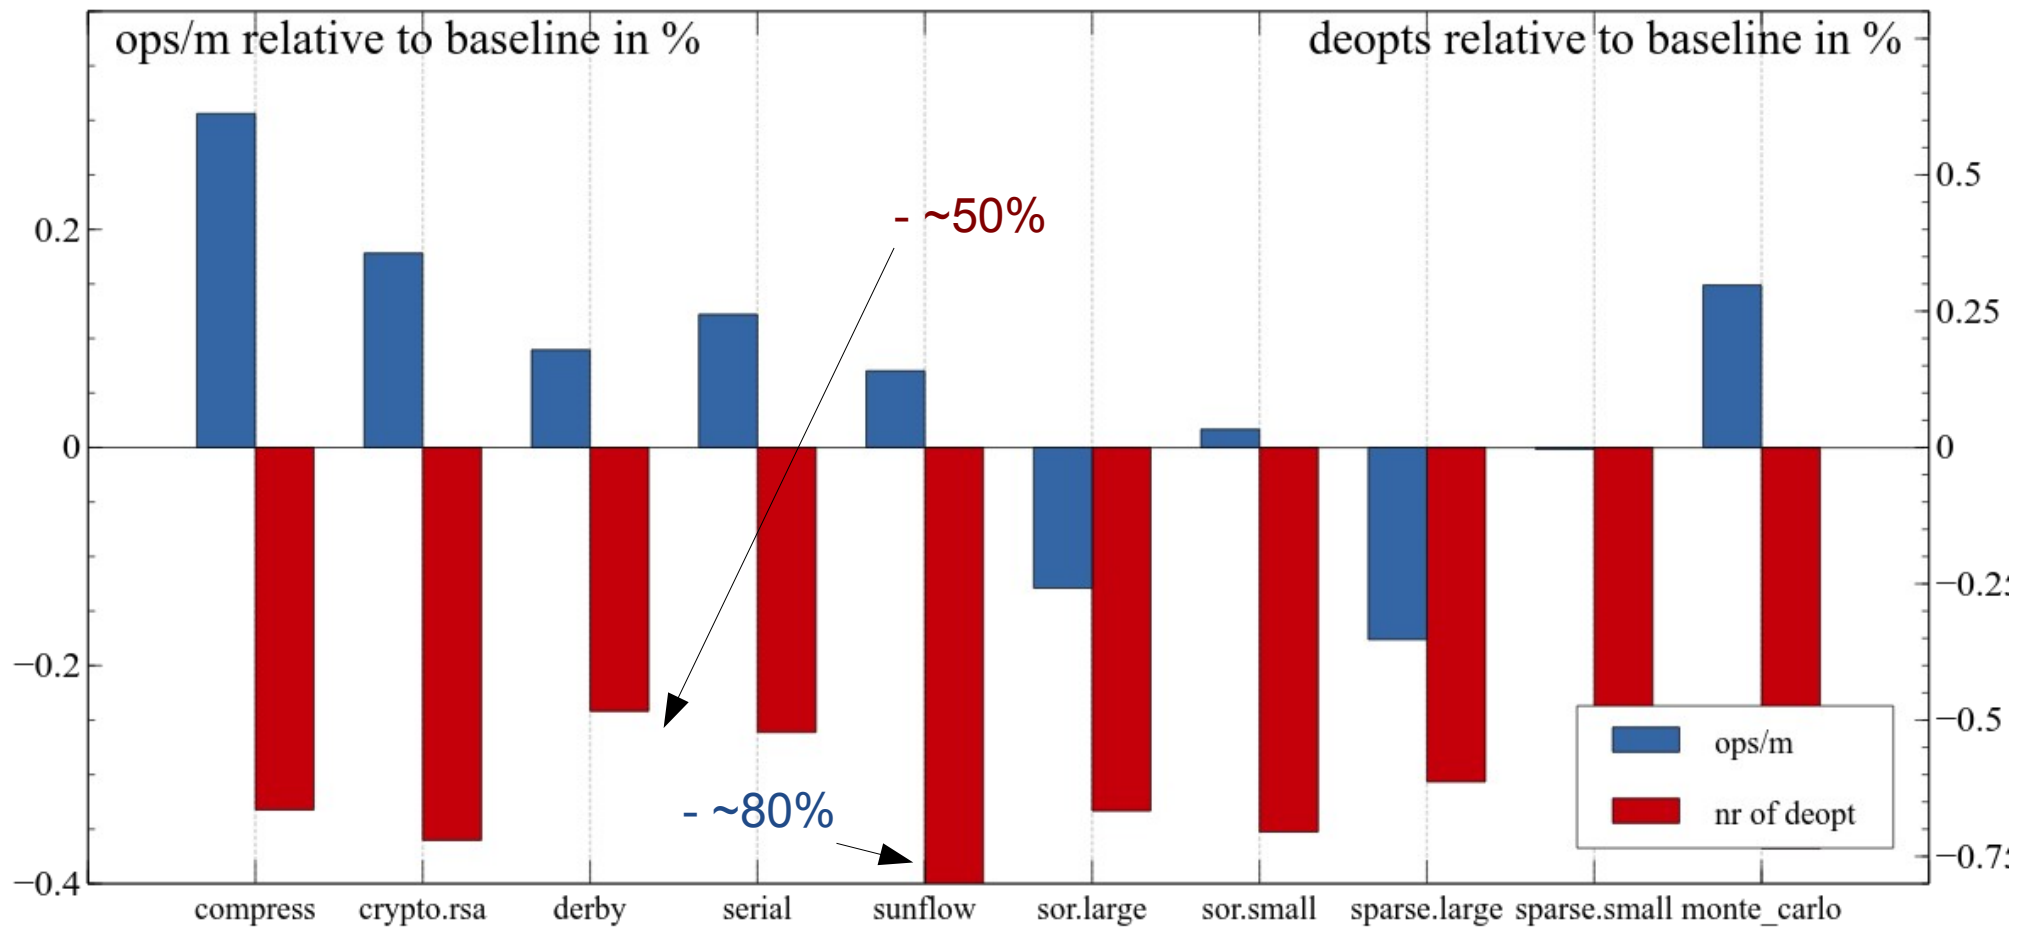
Performance evaluation

- **Performance** (higher is better)



Performance evaluation

- **Deoptimizations** (lower is better)

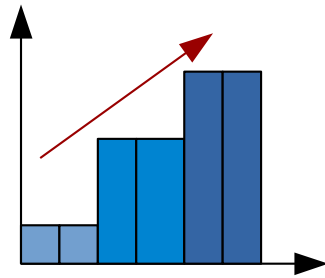


Other benchmark results


- Benefit mainly from cached **C2 compilations**.
Disabling interpreter profiles rarely affects performance
- Around 70% of the compilations **use cached profiles**
- No association between performance and load on **compile queue**

Conclusion

- Cached profiles can **improve** warmup performance
- System allows **fine tuning**
- **Main benefits:**



faster warmup




fewer deoptimizations

- Room for future work

End

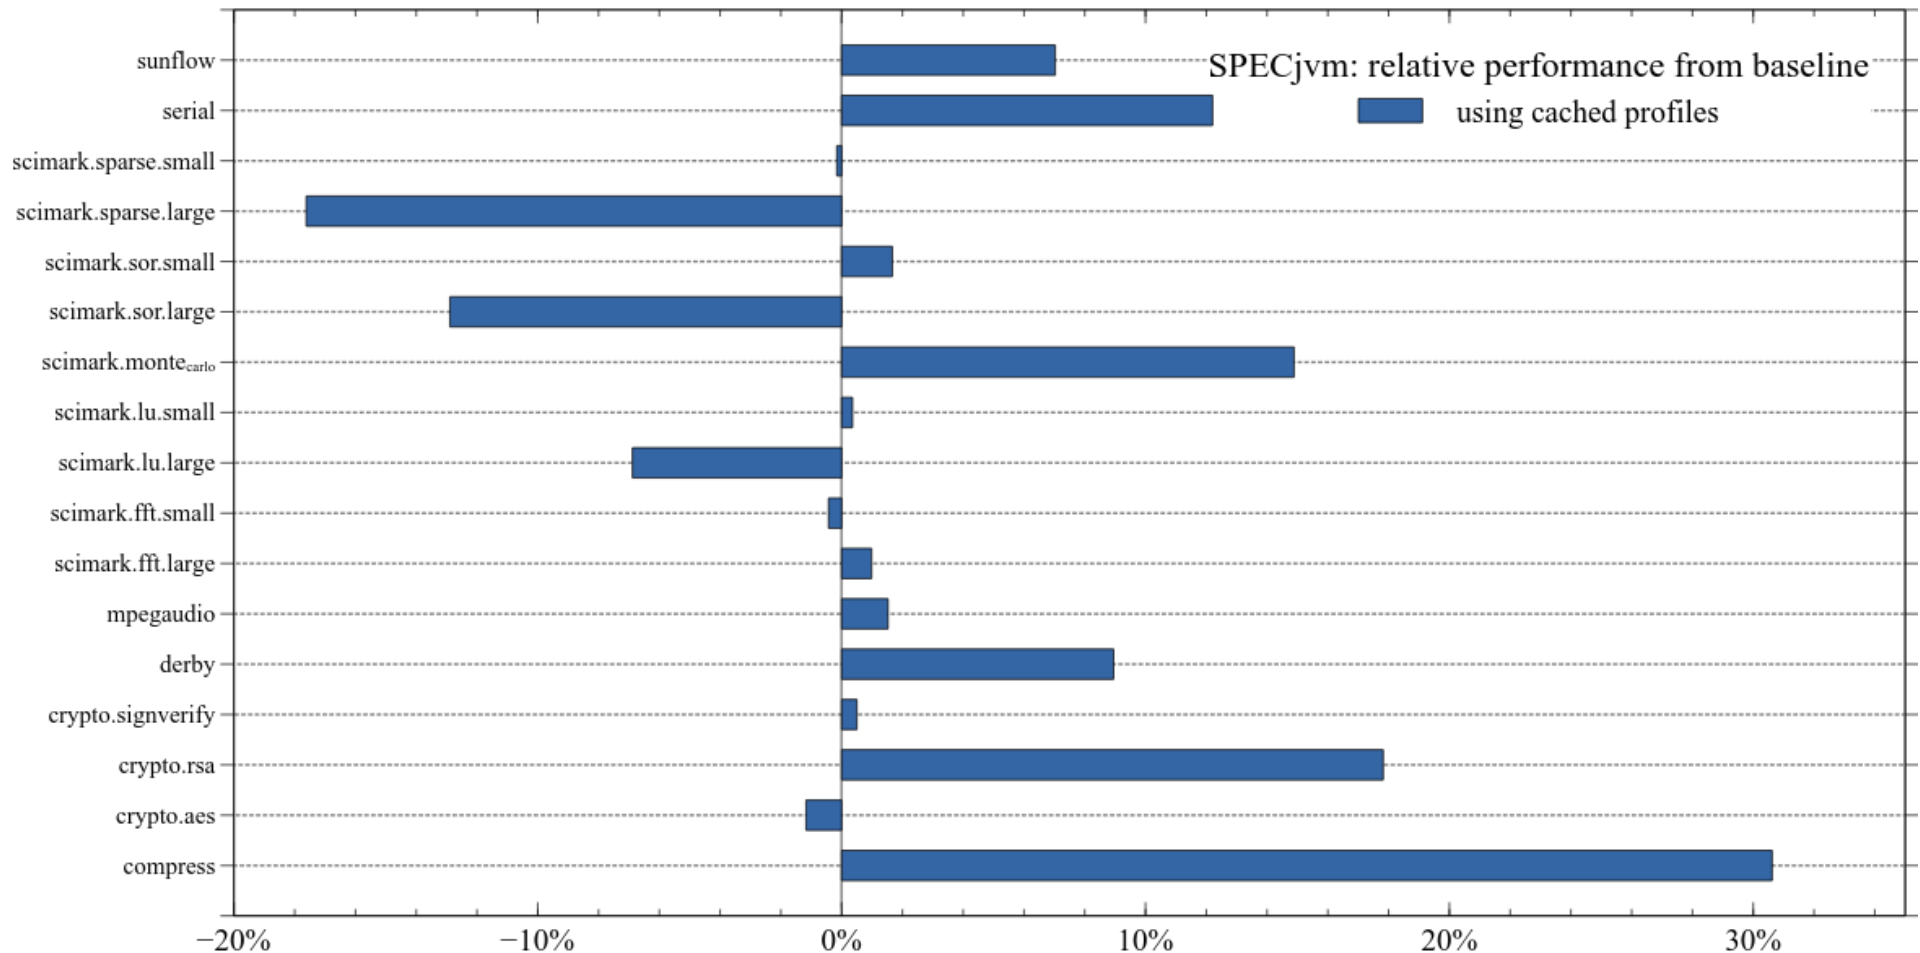
Additional graphes follow

SPECjvm: relative performance full runs (2+4 minutes)




SPECjvm: relative performance

warmup: all benchmarks




Octane: relative performance


all benchmarks



SPECjvm: relative performance without interpreter profiles




Octane: relative performance without interpreter profiles




Type of compilations

compress - mode 2



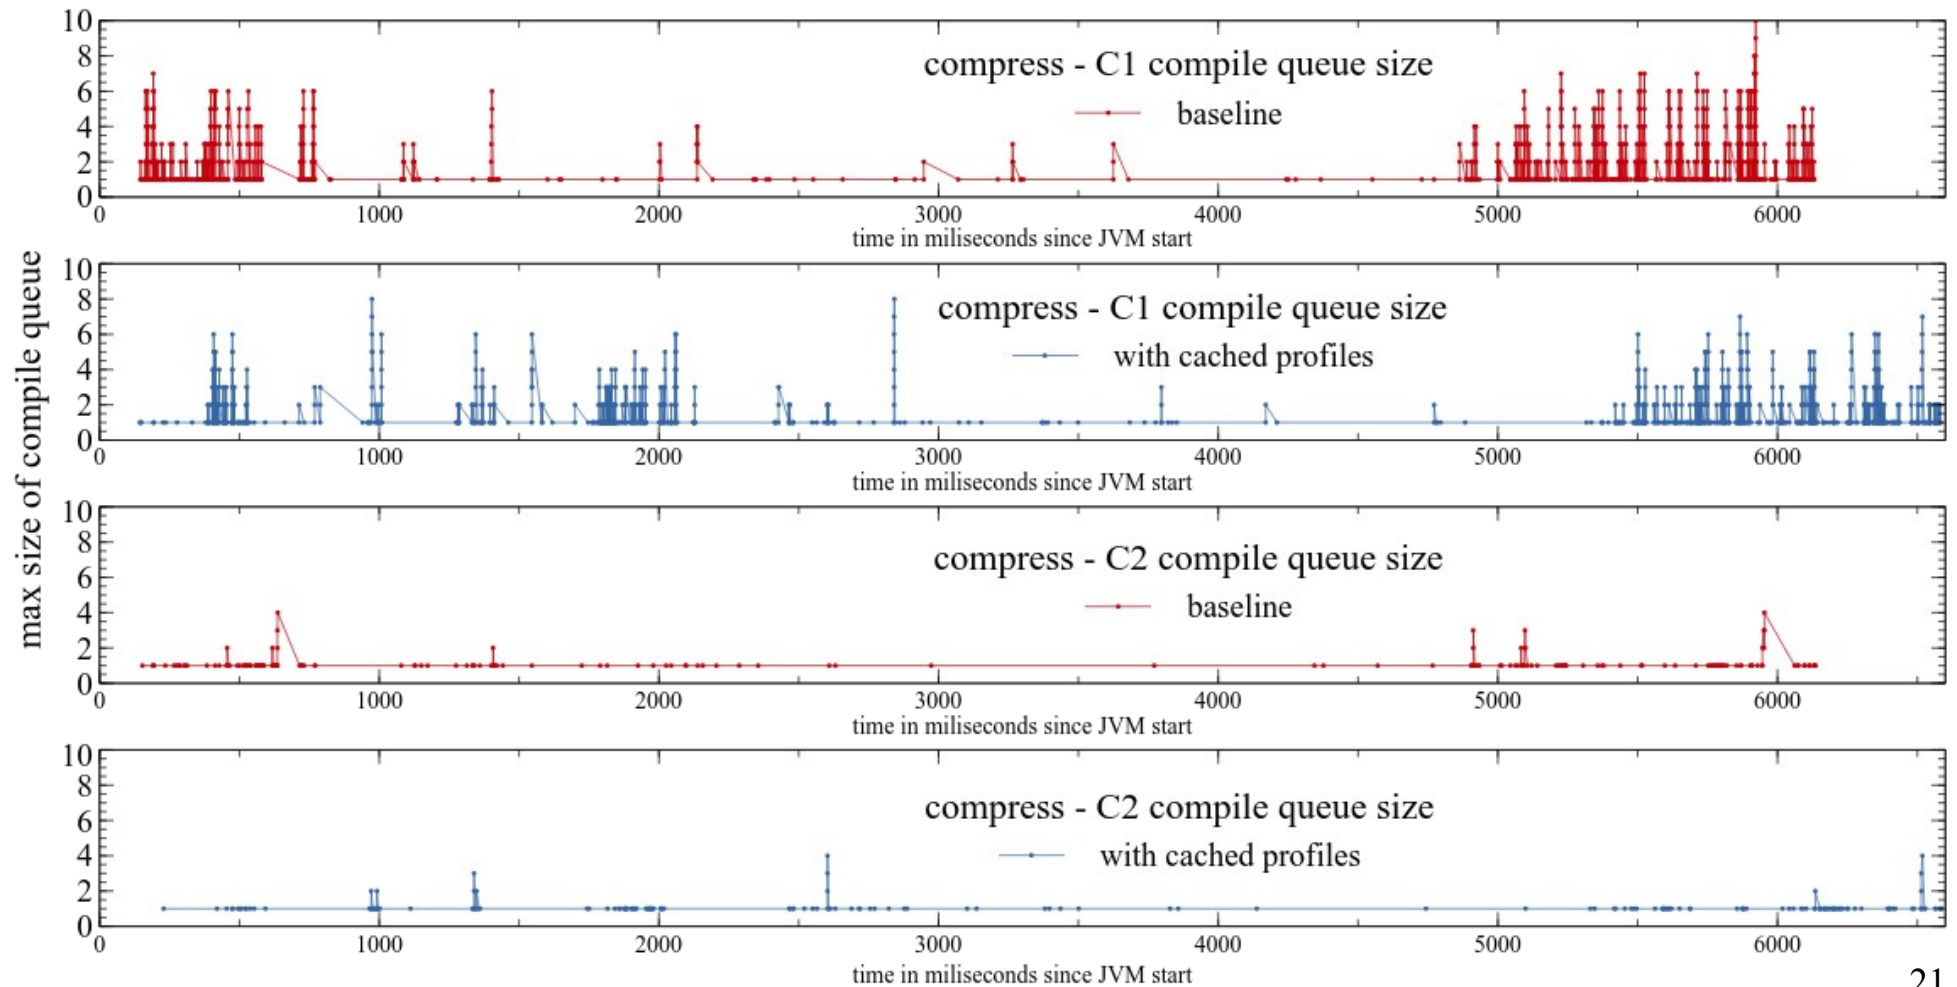
sparse.large - mode 2



- without cached profile – compile level 1
- without cached profile – compile level 2
- without cached profile – compile level 3
- without cached profile – compile level 4
- with cached profile – compile level 3
- with cached profile – compile level 4

Compile queue

- compress



Compile queue

- scimark.sparse.large

