

# Profile Caching for the Java Virtual Machine

Marcel Mohler, ETH Zurich  
Bachelor Thesis



*Supervisors: Zoltan Majo, Oracle  
Tobias Hartmann, Oracle*

*Prof. Thomas Gross, Laboratory for Software Technology*

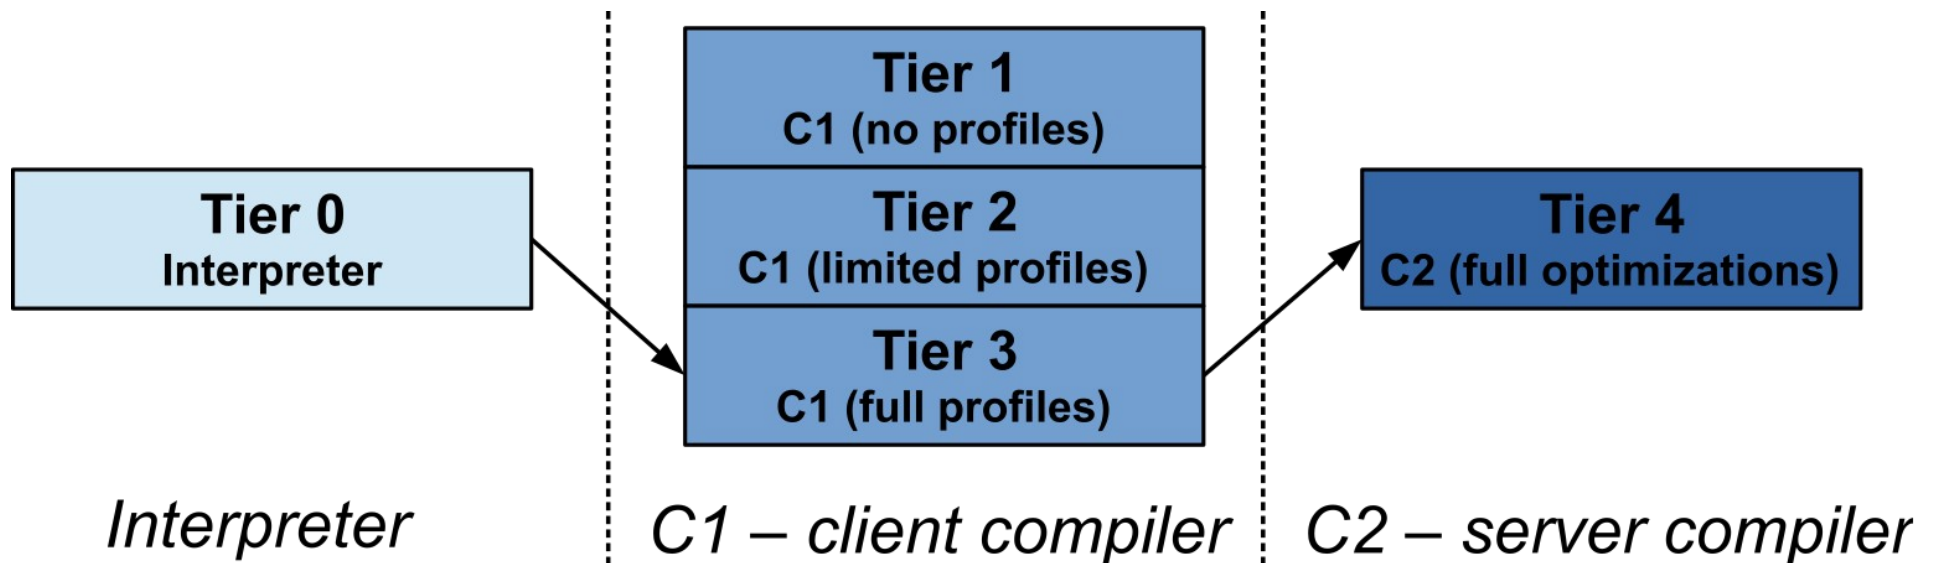
**ETH**

Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

**LST**

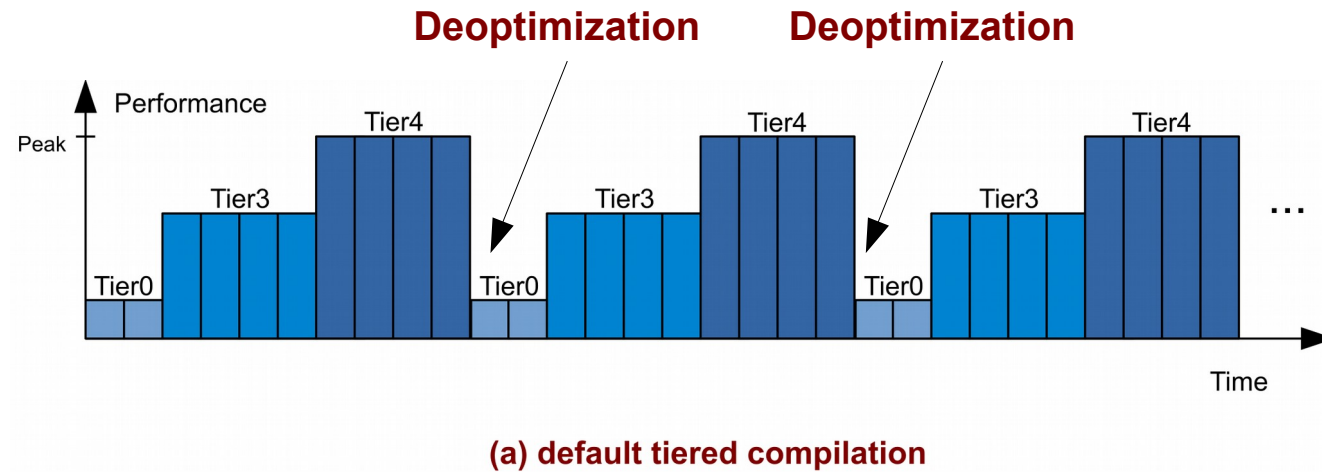
Laboratory for Software Technology

# Hotspot™: Tiered Compilation




- JVM gathers profiles
- Uses these profiles for code optimizations


# Problem



# Idea



(a) default tiered compilation





(b) with cached profiles

No deoptimization

# Goals

## Decrease performance fluctuations

- Faster method performance warmup  
→ reach peak performance quicker
- Less deoptimizations  
→ stay on peak performance



# Design: dump profiles

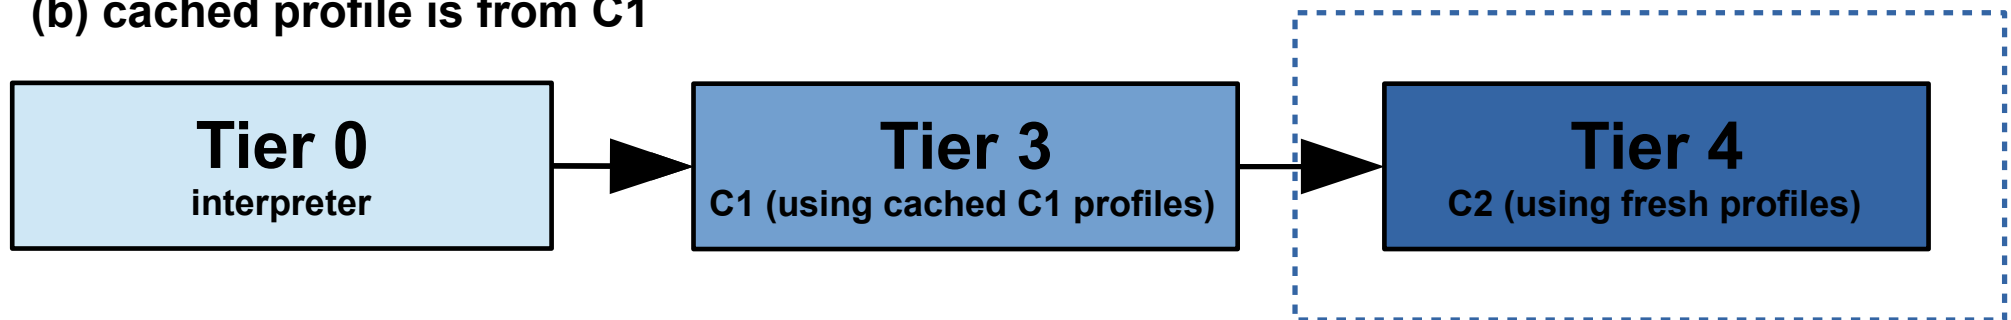
- 1 run of JVM where profiles get dumped to disk
- Store method metadata, profiles and compile information of C3 & C4 compilations

# Design: use profiles

(a) cached profile is from C2:



(b) cached profile is from C1



# Implementation

- 1846 lines of code
- 24 files affected
- 2 new classes
  - ciCacheProfiles
  - ciCacheProfilesBroker




# Evaluation

- ETH Data Center Observatory
- Focus on **warmup**, not overall performance
- 2 **benchmark suites**
  - SPECjvm 2008  
17 individual benchmarks
  - Google Octane (using Nashorn),  
16 individual benchmarks

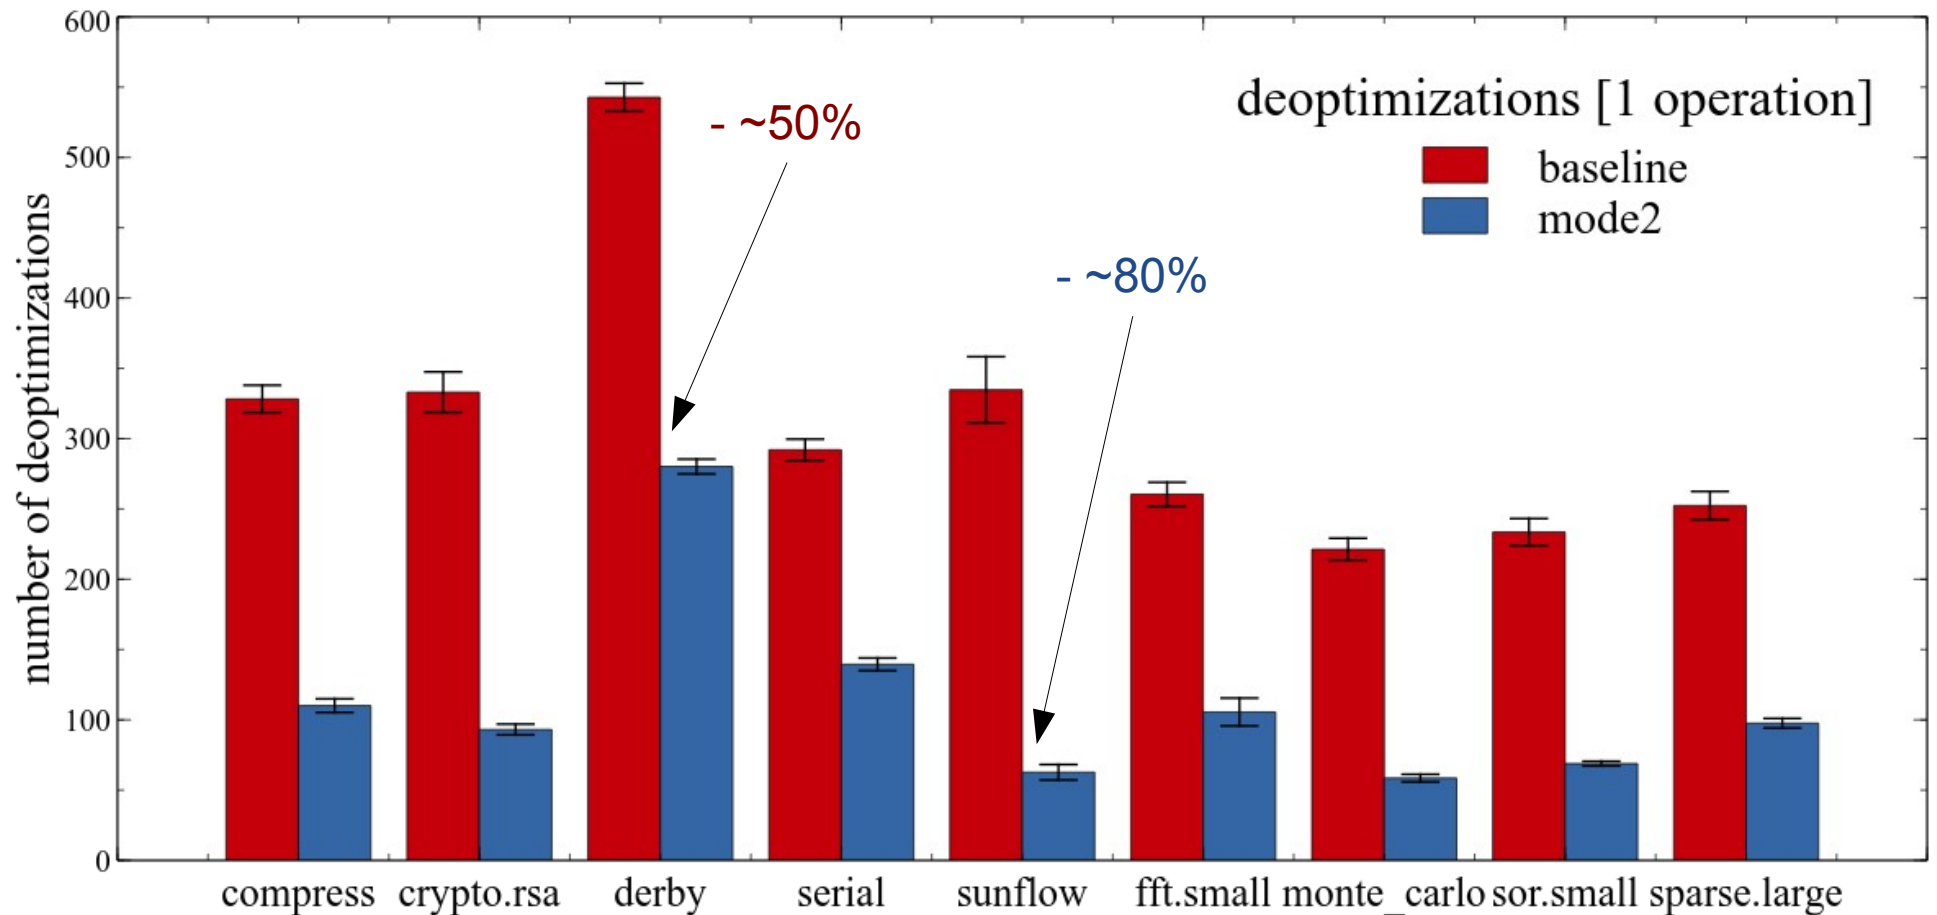
# Performance evaluation

- Performance (higher is better)



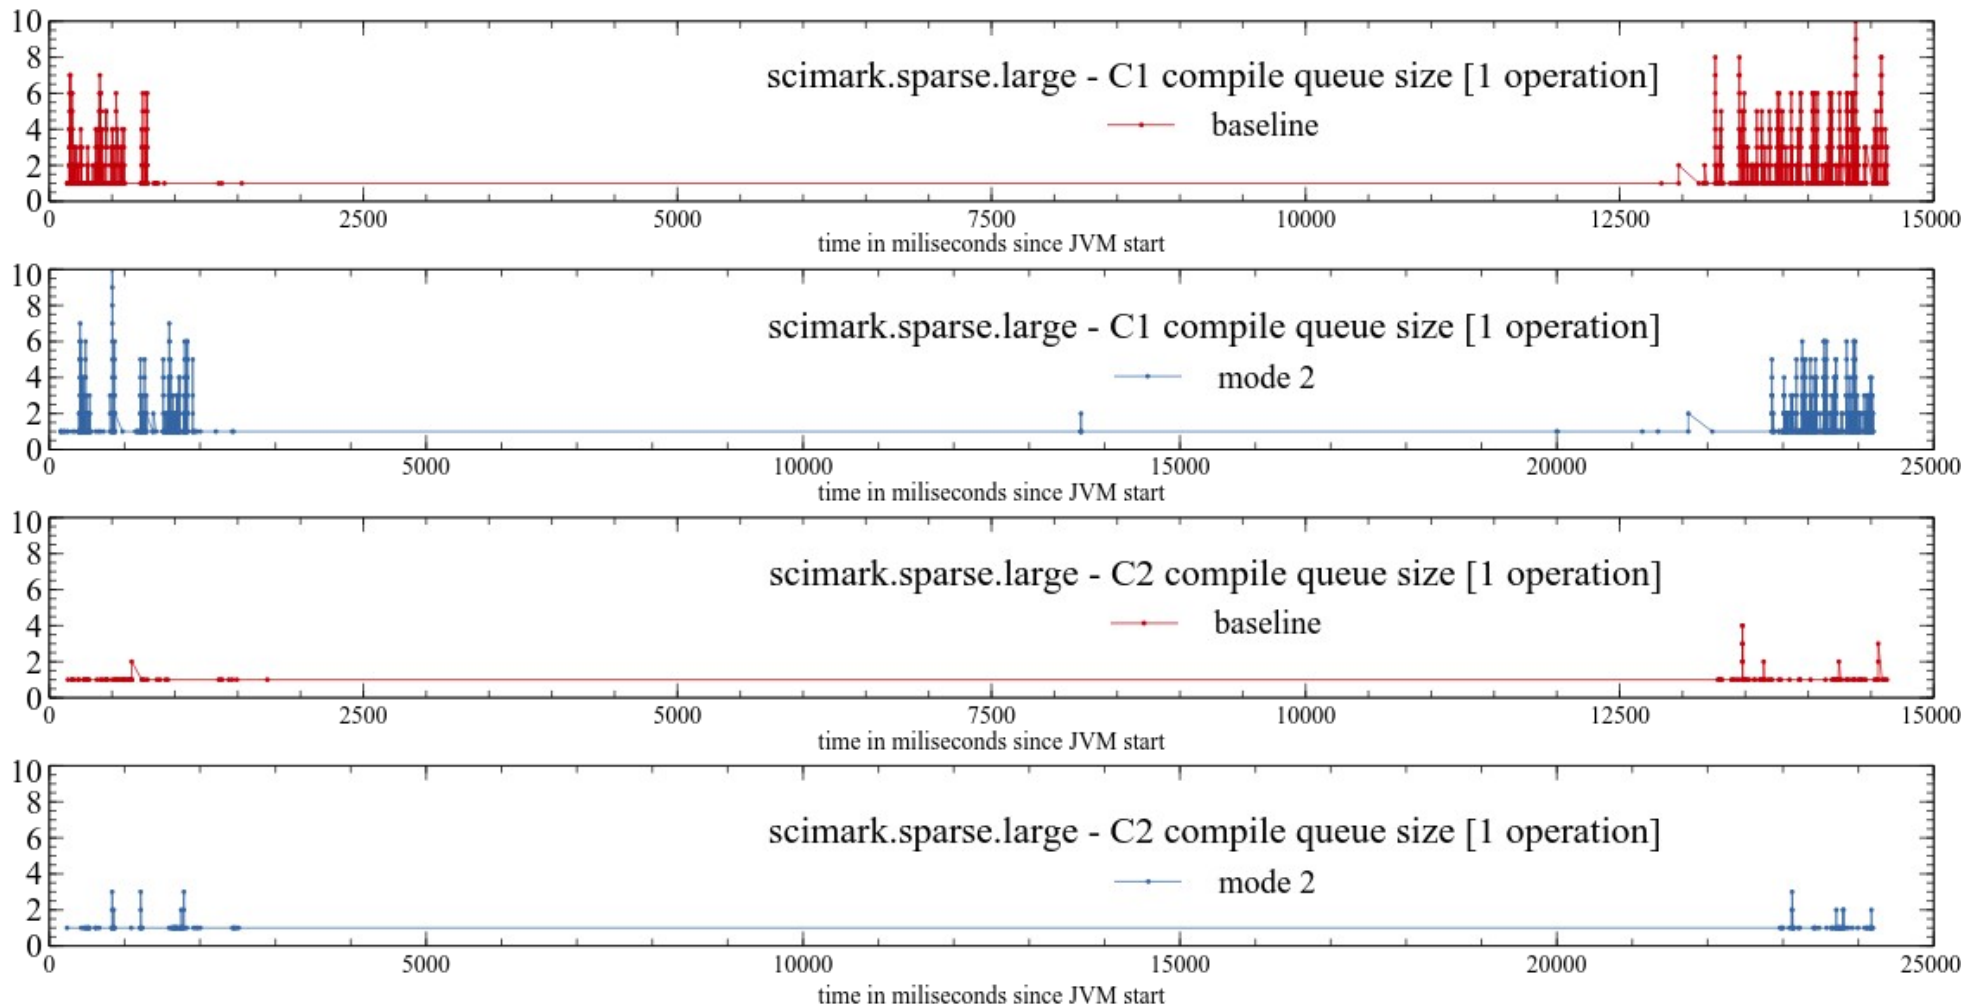
# Performance evaluation

- Deoptimizations (lower is better)



# Performance evaluation

- Compilation queue



# Other benchmark results

- Disabling **intrinsic**s does not influence performance
- Benefit mainly from **C2 compilations**. Disabling C1 profiles does not affect performance significantly
- Around 70% of the compilations **use profiles**

# Other approaches

- Presented: Mode2
- Mode 0: skip C1 & lower compilation thresholds
- Mode 1: skip C1 & keep original compilation thresholds

# Conclusion

- **Complex** system
- **Reasons** for performance influence **difficult** to measure
- Cached profiles *can* **greatly improve** warmup performance if used properly
- System requires **manual configuration**

Thank you for listening



# Profile Caching for the Java Virtual Machine

Marcel Mohler, ETH Zurich  
Bachelor Thesis



*Supervisors: Zoltan Majo, Oracle  
Tobias Hartmann, Oracle*

*Prof. Thomas Gross, Laboratory for Software Technology*

**ETH**

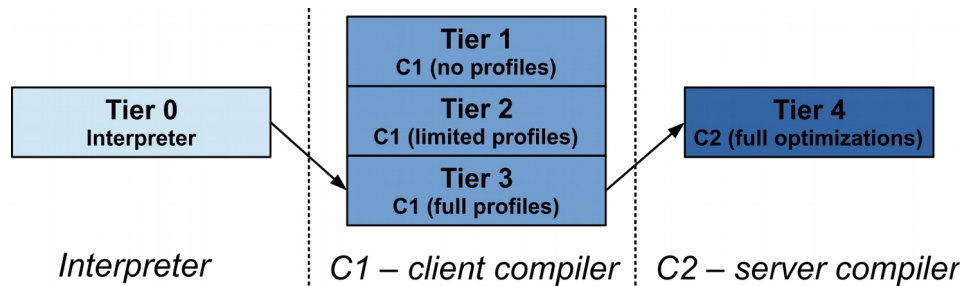
Eidgenössische Technische Hochschule Zürich  
Swiss Federal Institute of Technology Zurich

**LST**

Laboratory for Software Technology

## Welcome

## Hotspot™: Tiered Compilation




- JVM gathers profiles
- Uses these profiles for code optimizations

First off, a few things about the Just-in-time compiler of Hotspot

Hotspot is a JVM maintained by Oracle and heavily used all around the world in millions of devices

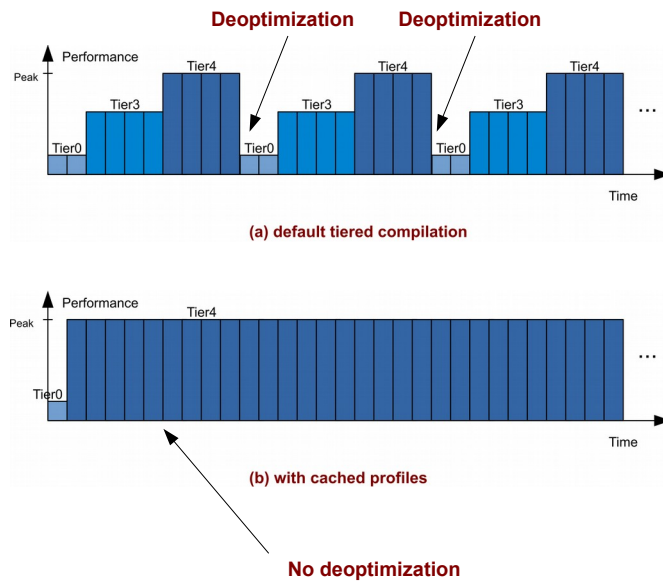
Profiles crucial to the performance of the code

# Problem



- What is the current problem?
- Methods start compiling at Tier0 and have to go through multiple tiers until reaching peak performance
- In case of deoptimizations we start from the beginning again
- Ideally we want this

# Idea




- Don't rely on the freshly gathered profiles
- Use profiles created in a previous run of the same program
- Allows the jvm to reach peak performance more quickly


# Goals

## Decrease performance fluctuations

- Faster method performance warmup  
→ reach peak performance quicker



- Less deoptimizations  
→ stay on peak performance



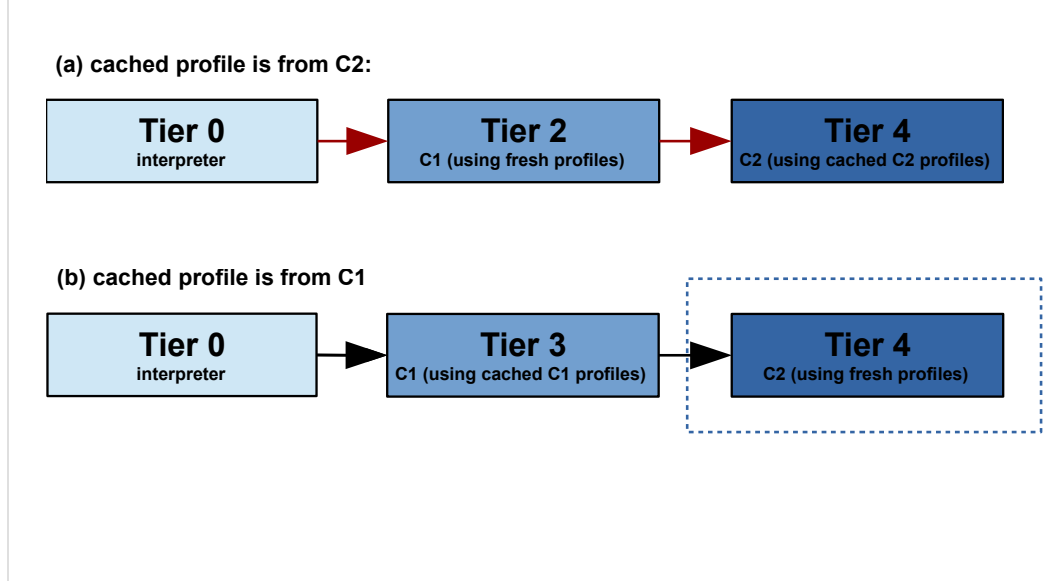
## Design: dump profiles

- 1 run of JVM where profiles get dumped to disk
- Store method metadata, profiles and compile information of C3 & C4 compilations

Takes two steps

First dump the profiles

## Design: use profiles



In a future run of the JVM we can then use these profiles

They get stored in a data-structure located in the C++ heap

And used whenever possible

I present mode 2 with focus on keeping the original tiered compilation steps (being conservative), there are 2 more modes

## Implementation

- 1846 lines of code
- 24 files affected
- 2 new classes
  - ciCacheProfiles
  - ciCacheProfilesBroker

More than 1800 lines of code

Two dozens of files affected

Separated from the rest, turning feature of results in JVM being unaffected by the changes



## Evaluation

- ETH Data Center Observatory
- Focus on [warmup](#), not overall performance
- 2 [benchmark suites](#)
  - SPECjvm 2008  
17 individual benchmarks
  - Google Octane (using Nashorn),  
16 individual benchmarks

DCO provided by ETH

Had to focus on warmup since it doesn't really affect performance for a long running benchmark

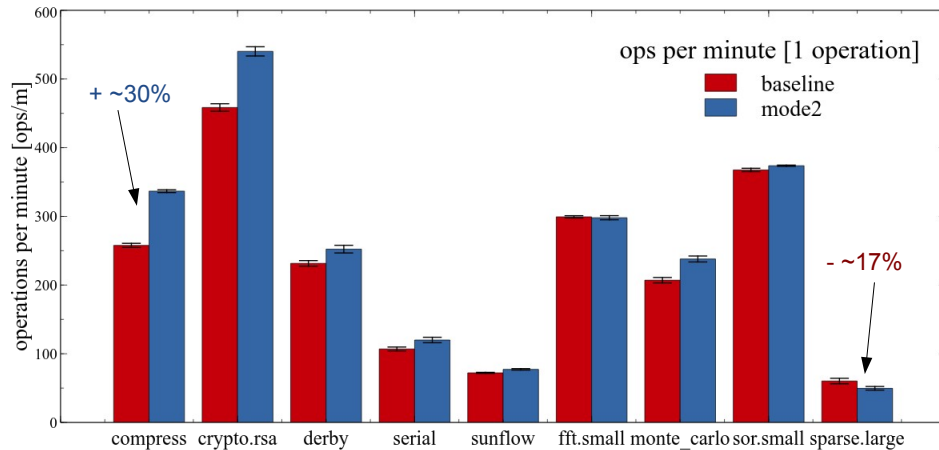
Divided benchmarks up in their parts and restarted JVM in between benchmarks

Also several microbenchmarks which are presented in the thesis

Going to present numbers from Mode2, there is a lot more in the thesis

# Performance evaluation

- **Performance** (higher is better)

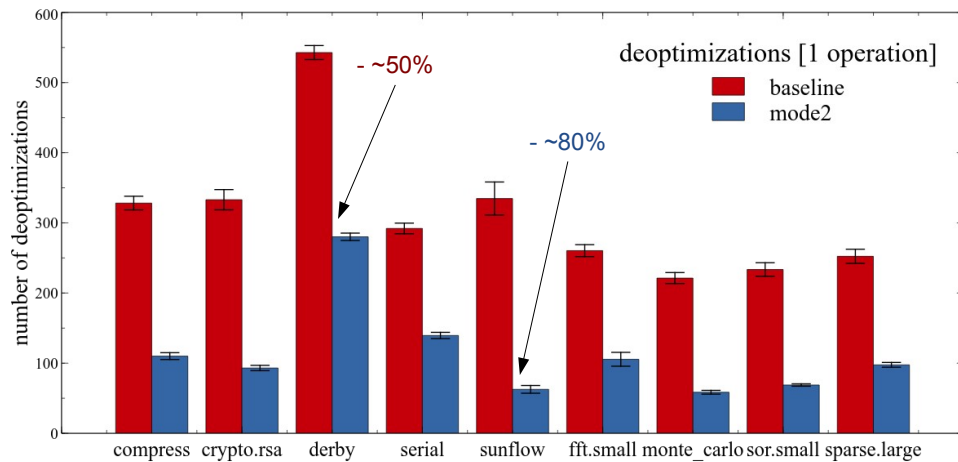


Performance example of SpecJVM benchmark

1 run, around 5-30 seconds

# Performance evaluation

- **Deoptimizations** (lower is better)

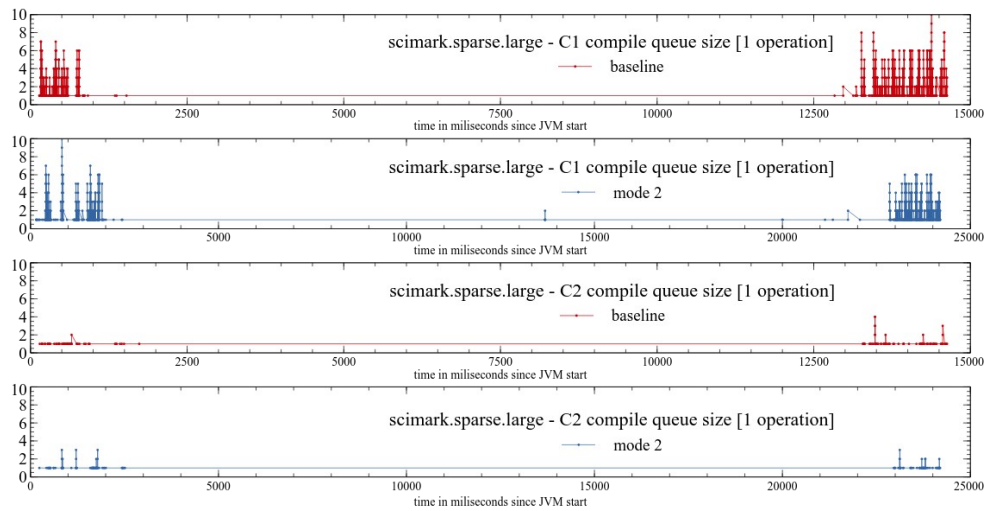


Deoptimizations lowered significantly, good indicator of improved code quality

Still difficult to measure how much influence a deopt has

# Performance evaluation

- Compilation queue



## Other benchmark results

- Disabling **intrinsic**s does not influence performance
- Benefit mainly from **C2 compilations**. Disabling C1 profiles does not affect performance significantly
- Around 70% of the compilations **use profiles**

Rest 30% are mainly L1/L2 compilations and a few skipped ones (lambda expressions etc)

## Other approaches

- Presented: Mode2
- Mode 0: skip C1 & lower compilation thresholds
- Mode 1: skip C1 & keep original compilation thresholds

## Conclusion

- **Complex** system
- **Reasons** for performance influence **difficult** to measure
- Cached profiles *can* **greatly improve** warmup performance if used properly
- System requires **manual configuration**

Thank you for listening