Profile Caching for the Java Virtual Machine

Marcel Mohler, ETH Zurich
Bachelor Thesis

Supervisors: Zoltan Majo, Oracle
            Tobias Hartmann, Oracle

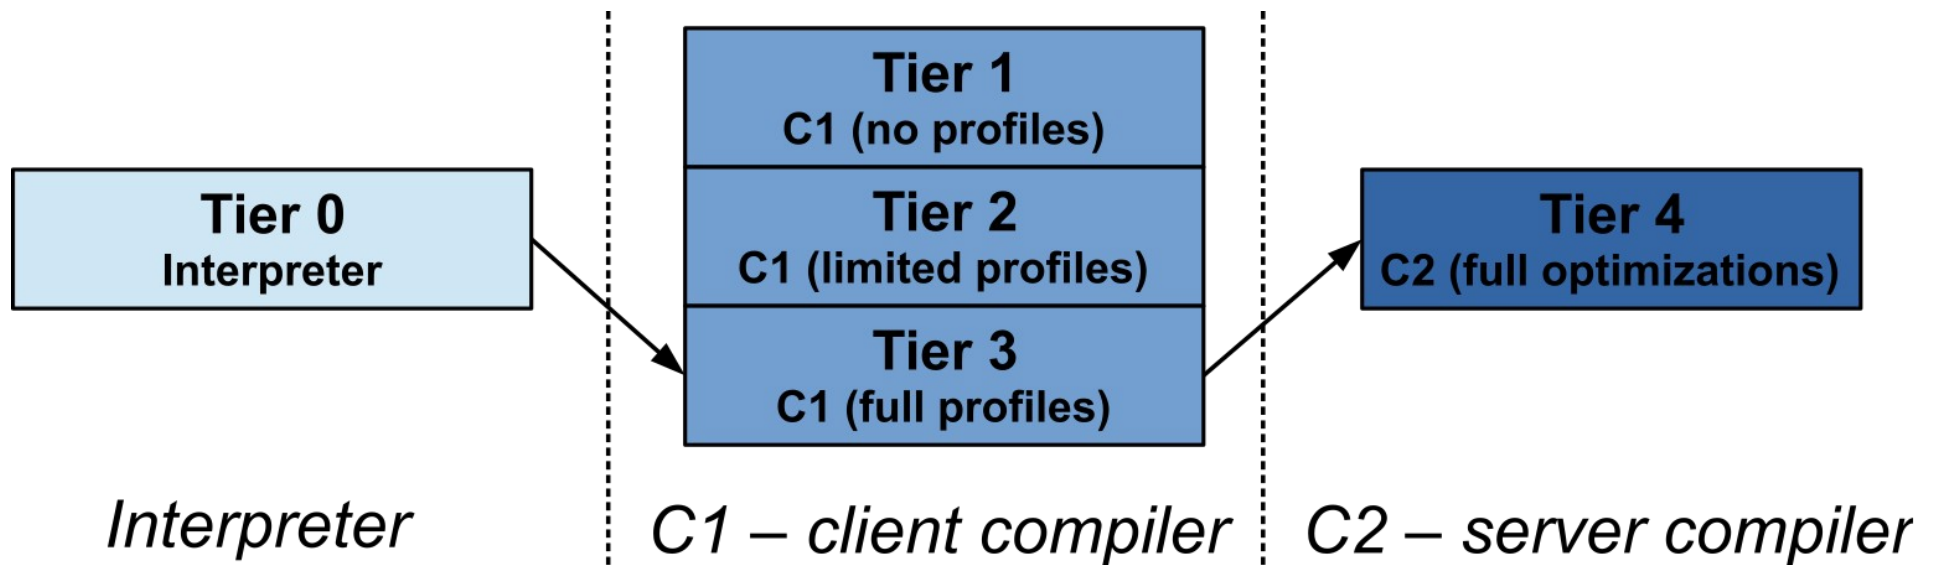*Prof. Thomas Gross, Laboratory for Software Technology*

ETH
Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

LST
Laboratory for Software Technology

Hotspot™: Tiered Compilation

- JVM gathers profiles
- Uses these profiles for code optimizations

Problem

(a) default tiered compilation

Idea

(a) default tiered compilation

(b) with cached profiles

Goals

Decrease performance fluctuations

- Faster method performance warmup

  → reach peak performance quicker

- Less deoptimizations

  → stay on peak performance

Design: dump profiles

- 1 run of JVM where profiles get dumped to disk

- Store method metadata, profiles and compile information of C3 & C4 compilations

Design: use profiles
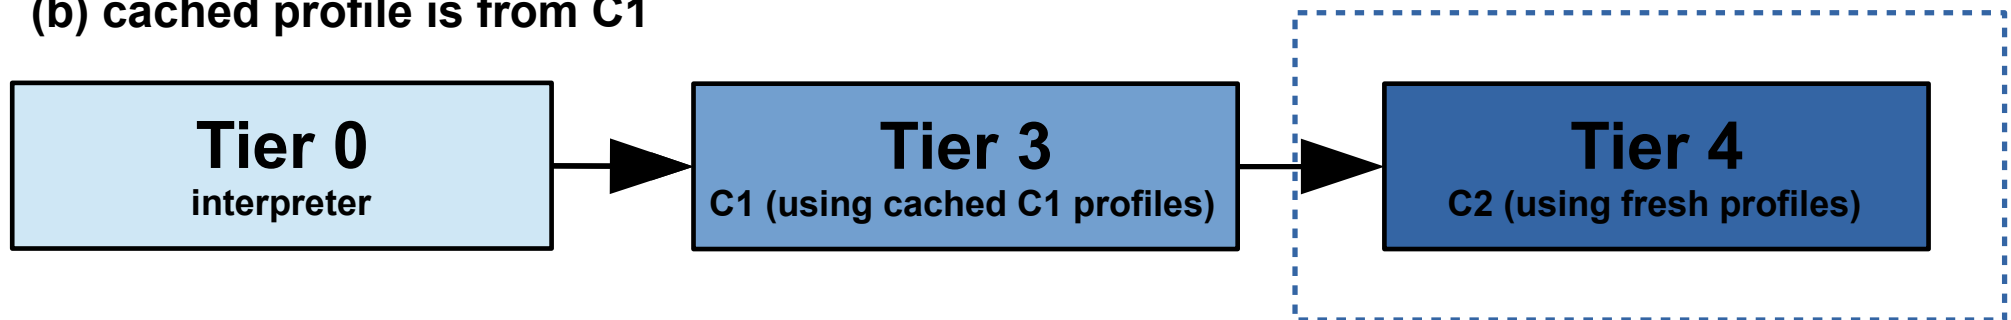
(a) cached profile is from C2:

Tier 0	Tier 2	Tier 4
interpreter	C1 (using fresh profiles)	C2 (using cached C2 profiles)

(b) cached profile is from C1

Tier 0	Tier 3	Tier 4
interpreter	C1 (using cached C1 profiles)	C2 (using fresh profiles)

Implementation

- 1846 lines of code

- 24 files affected

- 2 new classes
  - ciCacheProfiles
  - ciCacheProfilesBroker

Evaluation

- ETH Data Center Observatory

- Focus on warmup, not overall performance

- 2 benchmark suites

  - SPECjvm 2008

    17 individual benchmarks

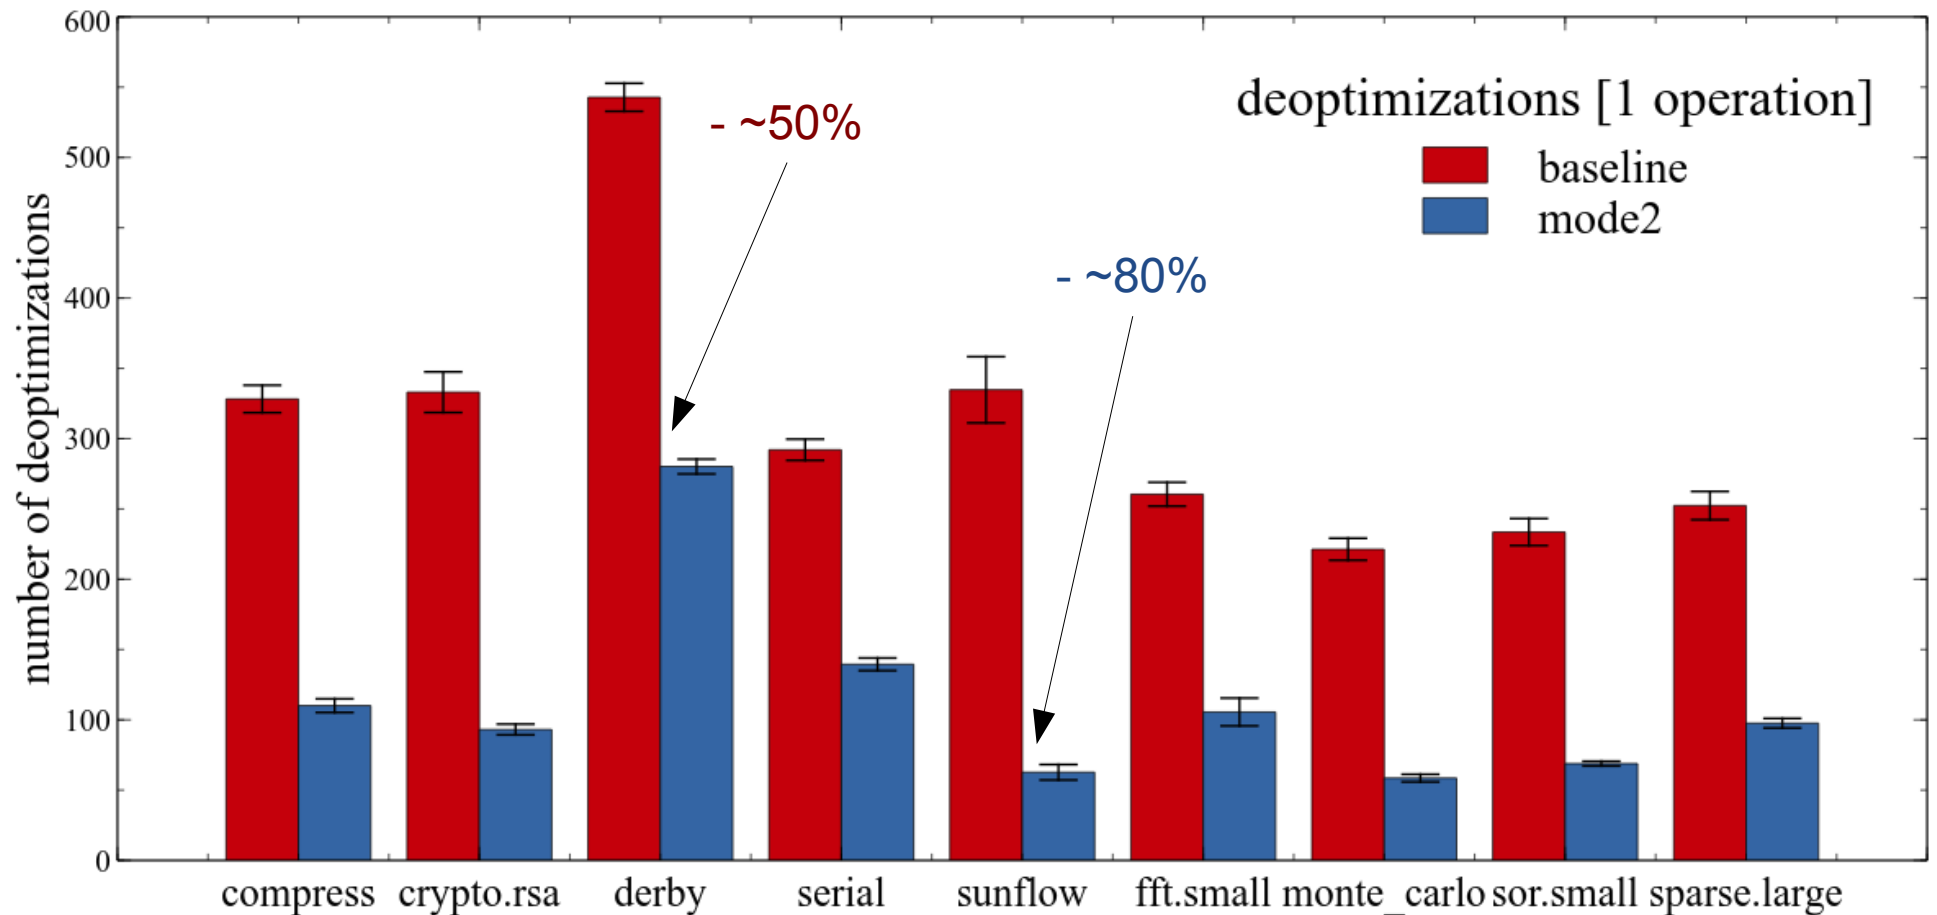 - Google Octane (using Nashorn)

    16 individual benchmarks

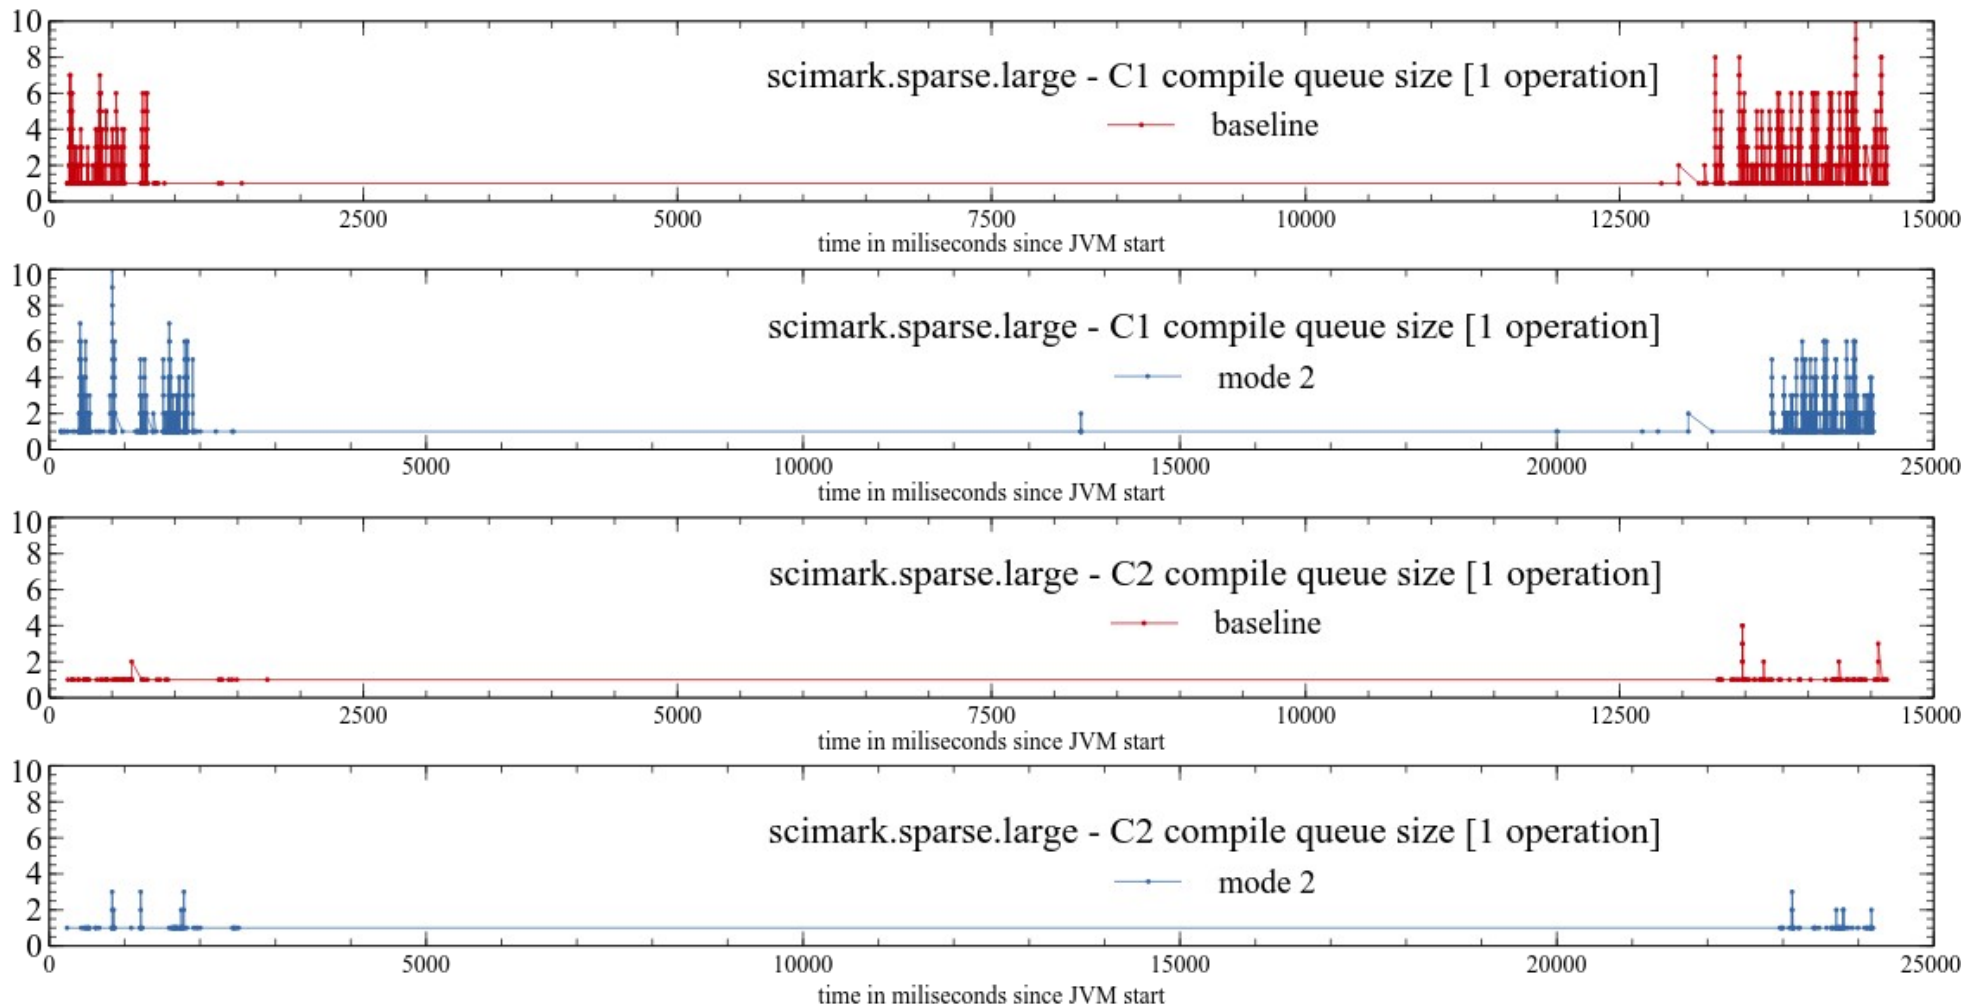# Performance evaluation

- Performance (higher is better)

Performance evaluation

- Deoptimizations (lower is better)

Performance evaluation

- Compilation queue

Other benchmark results

- Disabling intrinsics does not influence performance

- Benefit mainly from C2 compilations. Disabling C1 profiles does not affect performance significantly

- Around 70% of the compilations use profiles

Other approaches

- Presented: Mode2

- Mode 0: skip C1 & lower compilation thresholds

- Mode 1: skip C1 & keep original compilation thresholds

Conclusion

- Complex system

- Reasons for performance influence difficult to measure

- Cached profiles *can* greatly improve warmup performance if used properly

- System requires manual configuration

Thank you for listening

Profile Caching for the Java Virtual Machine

Marcel Mohler, ETH Zurich
Bachelor Thesis

Supervisors: Zoltan Majo, Oracle
Tobias Hartmann, Oracle

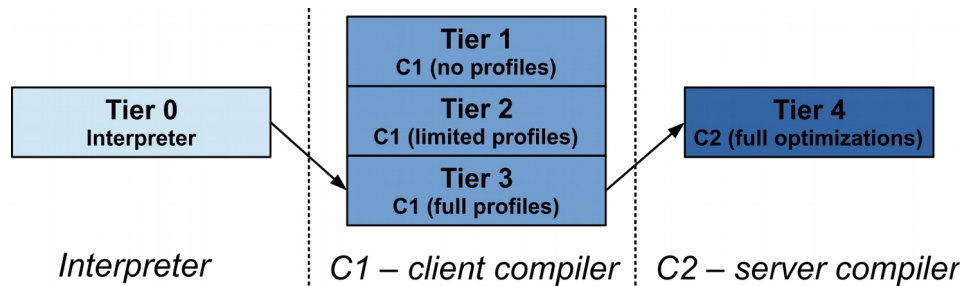*Prof. Thomas Gross, Laboratory for Software Technology*

ETH
Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

LST
Laboratory for Software Technology

Welcome

Hotspot™: Tiered Compilation

- JVM gathers profiles
- Uses these profiles for code optimizations

First off, a few things about the Just-in-time compiler of Hotspot

Hotspot is a JVM maintained by Oracle and heavily used all around the world in millions of devices

Profiles crucial to the performance of the code

Problem

(a) default tiered compilation

- What is the current problem?

- Methods start compiling at Tier0 and have to go through multiple tiers until reaching peak performance

- In case of deoptimizations we start from the beginning again

- Ideally we want this

Idea
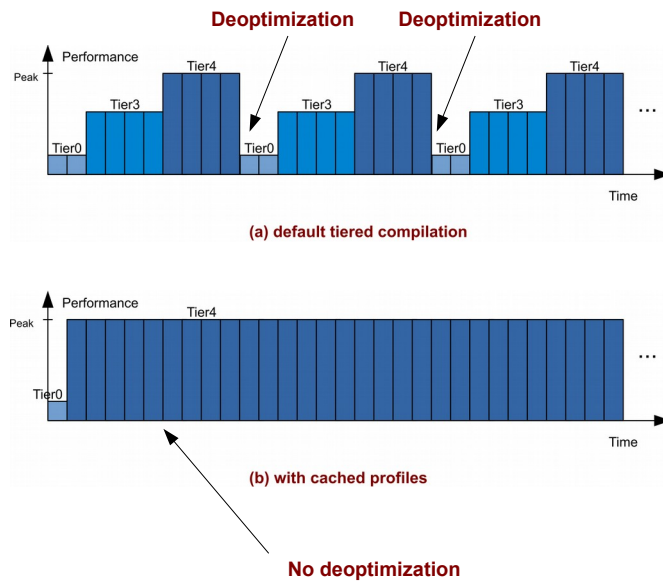
- Don't rely on the freshly gathered profiles
- Use profiles created in a previous run of the same program
- Allows the jvm to reach peak performance more quickly

Goals

Decrease performance fluctuations

- Faster method performance warmup
  - → reach peak performance quicker
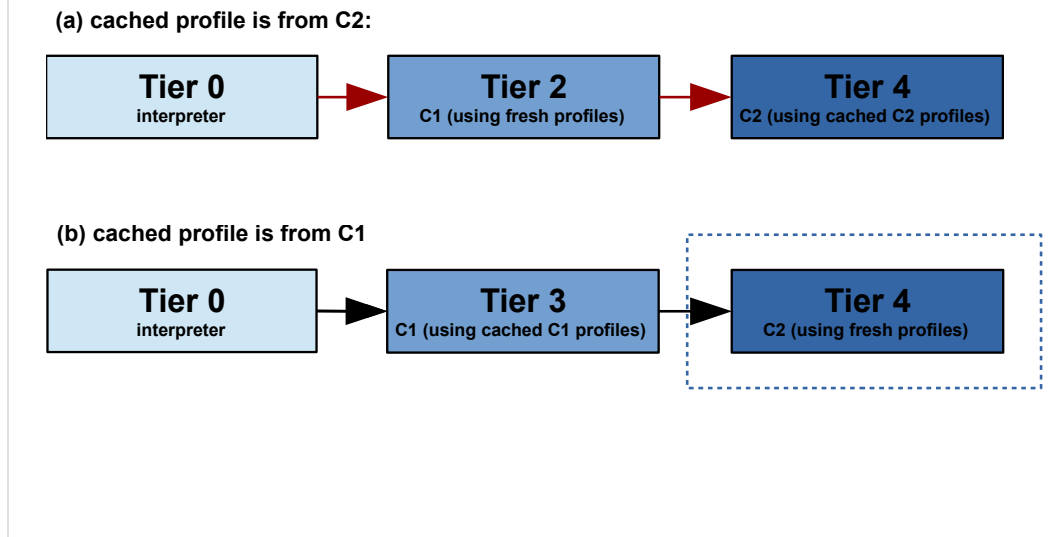
- Less deoptimizations
  - → stay on peak performance

Design: dump profiles

- 1 run of JVM where profiles get dumped to disk

- Store method metadata, profiles and compile information of C3 & C4 compilations

Takes two steps

First dump the profiles

Design: use profiles

(a) cached profile is from C2:

Tier 0	Tier 2	Tier 4
interpreter	C1 (using fresh profiles)	C2 (using cached C2 profiles)

(b) cached profile is from C1

Tier 0	Tier 3	Tier 4
interpreter	C1 (using cached C1 profiles)	C2 (using fresh profiles)

In a future run of the JVM we can then use these profiles

They get stored in a data-structure located in the C++ heap

And used whenever possible

I present mode 2 with focus on keeping the original tiered compilation steps (being conservative), there are 2 more modes

Implementation

- 1846 lines of code
- 24 files affected
- 2 new classes
  - ciCacheProfiles
  - ciCacheProfilesBroker

More than 1800 lines of code

Two dozens of files affected

Separated from the rest, turning feature of results in JVM being unaffected by the changes

Evaluation

- ETH Data Center Observatory
- Focus on warmup, not overall performance
- 2 benchmark suites
    - SPECjvm 2008
        17 individual benchmarks
    - Google Octane (using Nashorn)
        16 individual benchmarks

DCO provided by ETH

Had to focus on warmup since it doesn't really affect performance for a long running benchmark

Divided benchmarks up in their parts and restarted JVM in between benchmarks

Also several microbenchmarks which are presented in the thesis

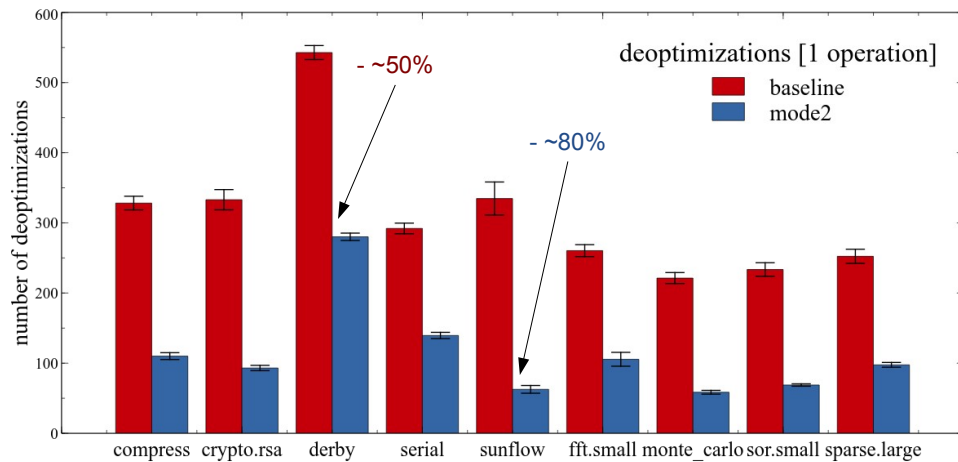Going to present numbers from Mode2, there is a lot more in the thesis

Performance example of SpecJVM benchmark

1 run, around 5-30 seconds

Performance evaluation
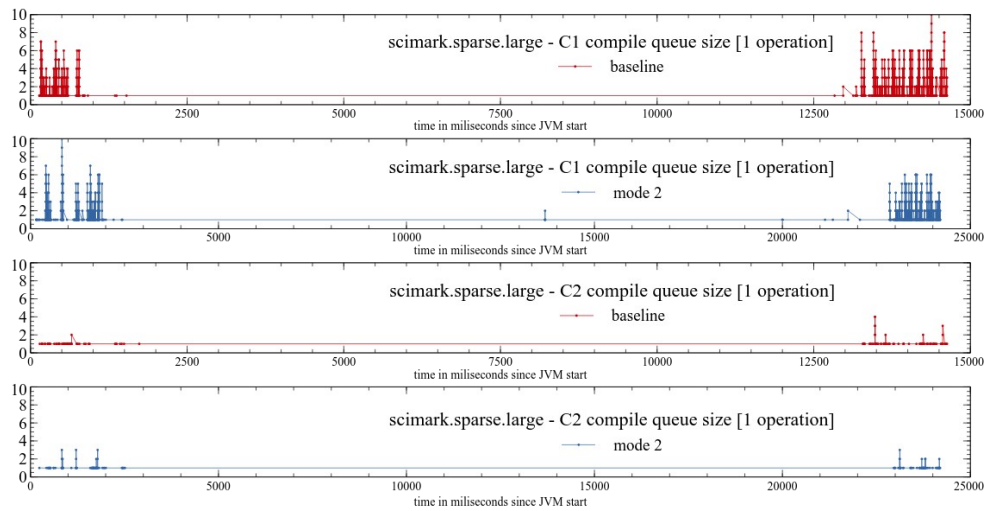
- Deoptimizations (lower is better)

Deoptimizations lowered significantly, good indicator of
  improved code quality

Still difficult to measure how much influence a deopt
  has

Performance evaluation

- Compilation queue

Other benchmark results

- Disabling intrinsics does not influence performance

- Benefit mainly from C2 compilations. Disabling C1 profiles does not affect performance significantly

- Around 70% of the compilations use profiles

Rest 30% are mainly L1/L2 compilations and a few skipped ones (lambda expressions etc)

Other approaches

- Presented: Mode2

- Mode 0: skip C1 & lower compilation thresholds
- Mode 1: skip C1 & keep original compilation thresholds

Conclusion

- Complex system
- Reasons for performance influence difficult to measure
- Cached profiles *can* greatly improve warmup performance if used properly
- System requires manual configuration

Thank you for listening