

Bachelor Thesis

Profile Caching for the Java Virtual Machine

Marcel Mohler

Zoltán Majó
Tobias Hartmann
Oracle Cooperation

Prof. Thomas R. Gross
Laboratory for Software Technology
ETH Zurich

August 2015



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich



Laboratory for Software Technology

Introduction

Virtual Machines like the Java Virtual Machine (JVM) are used as the execution environment of choice for many modern programming languages. The VMs interpret a suitable intermediate language (e.g., Java Byte Code for the JVM) and provide the runtime system for application programs and usually include a garbage collector, a thread scheduler, interfaces to the host operating system. As interpretation of intermediate code is time-consuming, VMs include usually a Just-in-Time (JIT) compiler that translates frequently-executed functions or methods to “native” code (e.g., x86 instructions).

The JIT compiler executes in parallel to a program’s interpretation by the VM, and as a result, compilation speed is a critical issue in the design of a JIT compiler. Unfortunately, it is difficult to design a compiler such that the compiler produces good (or excellent) code while limiting the resource demands of this compiler (the compiler requires storage and cycles – and even on a multi-core processor, compilation may slow down the execution of the application program). Consequently, most VMs adopt a multi-tier compilation system. At program startup, all methods are interpreted by the VM (execution at Tier 0). The interpreter performs profiling, and if a method is determined to be “hot”, this method is then compiled by the Tier 1 compiler. Methods compiled to Tier 1 are then profiled further and based on these profiling information, some methods are eventually compiled at Tier 2. One of the drawbacks of this setup is that for all programs, all methods start in Tier 0, with interpretation and profiling by the VM. However, for many programs the set of “hot” methods does not change from one execution to another and there is no reason to gather again and again the profiling information.

The main idea of this thesis is to cache these profiles from a prior execution to be used in further runs of the same program. This would allow the JIT compiler to use more sophisticated profiles early in program execution and avoid gathering the same profiling as well as prevent further compilations when more information about the method is available. While this will not influence the peak performance of the program, the hope is to decrease the time it’s needed to achieve it. I present an implementation on top of the Java Hotspot Virtual Machine as well as profound performance analysis using state-of-the-art benchmarks.

Contents

1	Overview Hotspot	1
1.1	Tiered Compilation	1
1.2	Deoptimizations	3
1.3	Compile Thresholds	3
1.4	On-Stack Replacement	4
2	Motivation	5
2.1	Example 1	6
2.2	Example 2	7
3	Implementation / Design	9
3.1	Creating cached profiles	9
3.2	Initializing cached profiles	10
3.3	Using cached profiles	11
3.4	Different usage modes for cached profiles	12
3.4.1	Compile Thresholds lowered (mode 0)	12
3.4.2	Unmodified Compile Thresholds (mode 1)	13
3.4.3	Modified C1 stage (mode 2)	13
3.5	Issues	14
3.6	Debug output	14
4	Performance	15
4.1	Setup	15
4.2	Startup performance	15
4.3	Deoptimizations	17
4.4	Effect on compile queue	17
5	Possible Improvements	21
6	Conclusion	23
A	Appendix	25
A.1	Tiered Compilation Thresholds	26
A.2	SPECjvm Benchmark	26

A.3 Octane Benchmark	28
Bibliography	30

1 Overview Hotspot

This chapter will provide the reader with an overview of the relevant parts of Java Hotspot. It explains the core concepts that are needed to understand the motivation and implementation of this thesis.

1.1 Tiered Compilation

As mentioned in the introduction, Programming Language Virtual Machines like Java Hotspot feature a multi-tier system when compiling methods during execution. Java VM's typically use Java Bytecode as input, a platform independent intermediate code generated by a Java Compiler like `javac` [7]. The Bytecode is meant to be interpreted by the virtual machine or further compiled into platform dependend machine code. Hotspot includes one interpreter and two different compilers with different profiling levels resulting in a total of 5 different tiers. Since in literature and the JVM source code use the *tiers* are also called *compilation levels* I use both as synonyms.

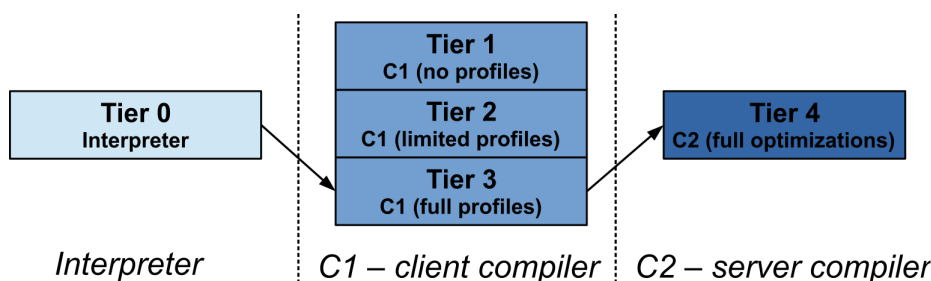


Figure 1.1: Overview of compilation tiers

All methods start being executed by Tier 0, also called the Interpreter. The interpreter is template-based, meaning for each bytecode instruction it emits a predefined assembly code snippet. During execution this code is also profiled. This means method information like execution counters, loop back-branches and additional statistics are counted. More importantly information about the program flow and state are gathered. These information contain for example which branches get taken or the final types of dynamically typed objects. Once one these counters exceed a predefined, constant threshold the method is considered *hot* which usually results in a compilation at a higher tier.

The standard behavior of Hotspot is to proceed with Level 3 (Tier 3). The method gets com-

piled with C1, also referred to as *client* compiler, and continues gathering full profiles. C1's goal is to provide a fast compilation with a low memory footprint. The client compiler performs simple optimizations such as constant folding, null check elimination and method inlining. It will also use the profiles gathered during interpretation. For example if certain branches were not taken during interpretation the C1 compiler might abstain from compiling these branches and provide a faster compilation.

The levels 1 and 2 include the same optimization but offer no or less profiling information and are used in special cases. Code compiled at these levels is significantly faster than level 3 because it needs to execute none or little instructions creating and managing the profiles. Anyway, since the profiles generated by C1 are further used in C2, Hotspot is usually interested in creating full profiles and therefor use level 3. There are however rare instances where a compilation of level 1 or level 2 is triggered. For example if enough profiles are available and a method can not be compiled by a higher tier, Hotspot might recompile the method with Tier 2 to get faster code until the higher tier compiler is available again. A compiler can become unavailable if its compilation queue exceeds a certain threshold.

More information about C1 can be found in [9] and [4].

Eventually, when further compile thresholds are exceeded, the JVM further compiles the method with C2, also known as *server* compiler. The server compiler makes use of the gathered profiles in Tier 0 and Tier 3 and produces highly optimized code. C2 includes far more and more complex optimizations like loop unrolling, common subexpression elimination and elimination of range and null checks. It performs optimistic method inlining, for example by converting some virtual calls to static calls. It relies heavily on the profiling information and richer profiles allow the compiler to use more and better optimizations. While the code quality of C2 is a lot better than C1 this comes at the cost of compile time. Since a C2 compilation includes A more detailed look at the server compiler can be found in [8]. Figure 1.1 gives a short overview as well as showing the standard transition.

The naming scheme *client/server* comes from back in the days where tiered compilation was not available and one had to choose the compiler via a Hotspot command line flag. The *client* compiler was meant to be used for interactive client programs with graphical user interfaces where response time is more important than peak performance. For long running server applications, the highly optimized but slower *server* compiler was used.

Tiered compilation was introduced to improve start-up performance of the JVM. Starting with the interpreter means that there is zero wait time until the method is executed since one does not need to wait until a compilation is finished. Also, consider that there are always parts of the code that get executed only once, where the compilation overhead would exceed the performance gain. C1 allows the JVM to have more optimized code available early which then can be used to

create a richer profile to be used when compiling with C2. Ideally this profile already contains most of the program flow and the assumptions made by C2 hold. If that is not the case the JVM might need to go back, gather more profiles and compile the method again. This is further described in the Section 1.2 *Deoptimizations*. In this case, being able to do quick compilations with C1 decreases the amount of C2 recompilations which are even more costly.

1.2 Deoptimizations

Ideally we compile a method with as much profiling information as possible. For example, since the profiling information are usually gathered in levels 0 and 3 it can happen that a method compiled by C2 wants to execute a branch it never used before. In this case the information about this branch are not available in the profile and therefore have not been compiled into the C2-compiled code. This is done to allow further, very optimistic optimization and to keep the compiled code smaller. So instead, the compiler places an uncommon trap at unused branches or unloaded classes which will get triggered in case they actually get used at a later time in execution.

The JVM then stops execution of that method and returns the control back to the interpreter. This process is called *deoptimization* and considered very expensive. The previous interpreter state has to be restored and the method will be executed using the slow interpreter. Eventually the method might get recompiled with the newly gained information.

1.3 Compile Thresholds

The transitions between the compilation levels (see Fig. 1.1) are chosen based on predefined constants called *compile thresholds*. When running an instance of the JVM one can specify them manually or use the ones provided. A list of thresholds and their default values relevant to this thesis are given in Appendix A.1. The standard transitions from Level 0 to 3 and 3 to 4 happen when the following predicate returns true:

$$i > TierXInvocationThreshold * s \\ || (i > TierXMinInvocationThreshold * s \ \&\& \ i + b > TierXCompileThreshold * s)$$

where X is the next compile level (3 or 4), i the number of method invocations, b the number of backedges and s a scaling coefficient (default = 1). The thresholds are relative and individual for interpreter and compiler.

On-stack replacement uses a simpler predicate:

$$b > TierXBackEdgeThreshold * s$$

Please note that there are further conditions influencing the compilation like the load on the compiler which will not be discussed.

1.4 On-Stack Replacement

Since the JVM does not only count method invocations but also loop back branches (see also Section 1.3) it can happen that a method gets compiled while it is still running and the compiled method is ready before the method has finished. Instead of waiting for the next method invocation Hotspot can replace the method directly on the program stack. This process is called *on-stack replacement* and usually shortened to OSR. The Figure 1.2 presented in a talk by T. Rodriguez and K. Russel [9] gives a graphical representation. The benefits of OSR will become more obvious when looking at the first example in Chapter 2.

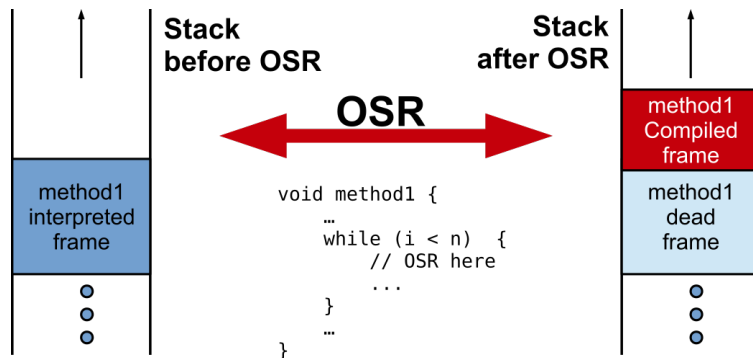


Figure 1.2: Graphical schema of OSR

2 Motivation

I continue with presenting two very simple example methods that illustrate the motivation and benefit from using cached profiles. This should provide the reader with an understanding how and why cached profiles can be beneficial for the performance of a Java Virtual Machine. I will omit any implementation details on purpose as they will be discussed in Chapter 3 in detail.

Ideally, being able to reuse the profiles from previous runs should result in two main advantages:

1. **Lower start-up time of the JVM:** Having information about the program flow already, the compiler can avoid gathering profiles and compile methods earlier and directly at higher compilation levels.
2. **Less Deoptimizations:** Since cache profiles get dumped at the end of a compilation, when using these profiles the compiler can already include all optimizations for all different method executions. Less uncommon traps need to be placed and less deoptimizations occur.




Figure 2.1: schematic visualization of cached profile benefit

Figure 2.1 gives a schematic visualization of the expected effect on performance of a single method when using cached profiles compared to the current state without such a system and standard tiered compilation. The blue bars roughly represent method invocations and higher bars equal higher compilation levels and therefore higher performance. The x-axis represents time since the start of the JVM. The figure shows the ideal case and abstracts away many details and other possible cases. However, it provides a good visualization for the examples provided in this chapter. A more detailed performance analysis, also considering worse cases is done in Chapter 4.

Listing 2.1: Simple method that does not get compiled

```

1 class NoCompile {
2     double result = 0.0;
3     for(int c = 0; c < 100; c++) {
4         result = method1(result);
5     }
6     public static double method1(double count) {
7         for(int k = 0; k < 10000000; k++) {
8             count = count + 50000;
9         }
10        return count;
11    }
12 }

```

I'm using my implementation described in Chapter 3 in `CachedProfileMode 0` (see 3.4.1) built into openJDK 1.9.0. All measurements in this chapter are done on a Dual-Core machine running at 2 GHz with 8GB of RAM. To measure the method invocation time I use `hprof` [6] and the average of 10 runs. The evaluation process has been automated using a couple of python scripts. The error bars show the 95% confidence interval.

2.1 Example 1

For this very first example, on-stack replacement has been disabled to keep the system simple and easy to understand.

Example one is a simple class that invokes a method one hundred times. The method itself consists of a long running loop. The source code is shown in Listing 2.1. Since OSR is disabled and a compilation to level 3 is triggered after 200 invocations this method never leaves the interpreter. I call this run the *Baseline*. To show the influence of cached profiles I use a compiler flag to lower the compile threshold explicitly and, using the functionality written for this thesis, tell Hotspot to cache the profile. In a next execution I use these profiles and achieve significantly better performance as one can see in Figure 2.2. This increase comes mainly from the fact that having a cached profile available allows the JVM to compile highly optimized code for hot methods earlier (at a lower threshold) since there is no need to gather the profiling information first.

Since the example is rather simple neither the baseline nor the profile usage run trigger any deoptimization. This makes sense because after the first invocation, all the code paths of the method have been taken already and are therefore known to the interpreter and saved in the profile.

Enabling OSR again and the difference between with and without cached profiles vanishes. This happens because Hotspot quickly realizes the hotness of the method and the JIT compiler produces perfectly optimized code during the first method invocation already. Even the OSR compiled code never triggers any deoptimization due to the simplicity of the loop. So this example appears rather artificial since the same performance can be achieved with OSR already but nevertheless shows the influence of early compilation.




Figure 2.2: NoCompile.method1 - per method invocation time

2.2 Example 2

However, OSR is one of the main features of Hotspot to improve the JIT performance and disabling that does not give us any practical results. Since we want an example which demonstrates the influence of cached profiles, I came up with the example sketched in Listing 2.2 which is slightly more complex but still easy to understand.

The idea is to create a method that takes a different, long running branch on each of its method invocations. Each branch has been constructed in a way that it will trigger an OSR compilation. When compiling this method during its first iteration only the first branch will be included in the compiled code. The same will happen for each of the 100 method invocations. As one can see in Figure 2.3 the baseline indeed averages at around 130 deoptimizations and a time per method invocation of 200 ms.

Now I use a regular execution to dump the profiles and then use these profiles. So theoretically the profiles dumped after a full execution should include knowledge of all branches and therefore the compiled method using these profiles should not run into any deoptimizations. As one can see in Figure 2.3 this is indeed the case. When using the cached profiles no more deoptimizations occur and because less time is spent profiling and compiling the methods the per method execution time is even significantly faster with averaging at 190ms now.

Listing 2.2: Simple method that causes many deoptimizations

```

1 class ManyDeopts {
2     double result = 0.0;
3     for(int c = 0; c < 100; c++) {
4         result = method1(result);
5     }
6     public static long method1(long count) {
7         for(int k = 01; k < 100000001; k++) {
8             if (count < 100000001) {
9                 count = count + 1;
10            } else if (count < 300000001) {
11                count = count + 2;
12            }
13            .
14            .
15            } else if (count < 505000000001) {
16                count = count + 100;
17            }
18            count = count + 50000;
19        }
20        return count;
21    }
22 }

```

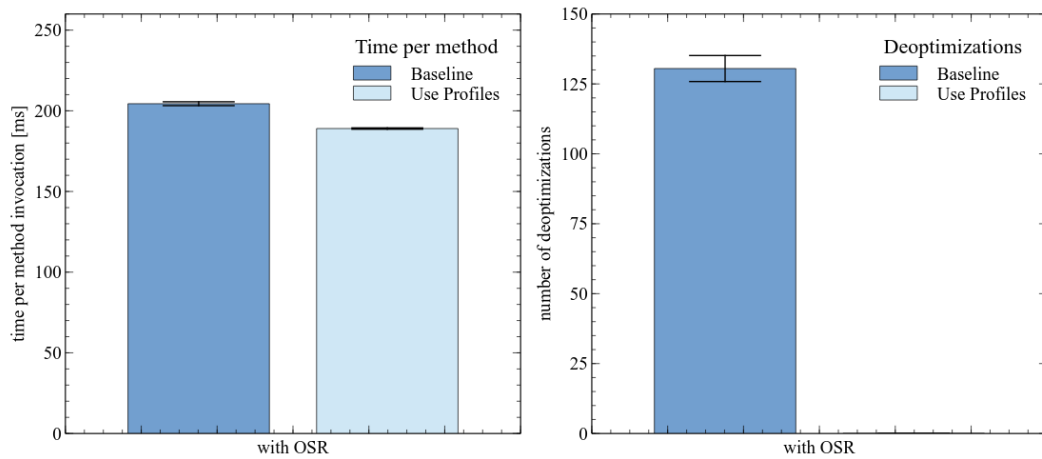


Figure 2.3: ManyDeopts.method1 - per method invocation time and deoptimization count

3 Implementation / Design

This chapter describes the implementation of the cached profiles implementation for Hotspot, written as part of this thesis.

Hotspot is an Java virtual machine implementation maintained by Oracle Corporation. It is part of the open source project `OpenJDK` and the source code is available at <http://openjdk.java.net/>.

Most of the work is included in two new classes `/share/vm/ci/ciCacheProfiles.cpp` and `/share/vm/ci/ciCacheProfilesBroker.cpp` as well as modifications to `/share/vm/ci/ciEnv.cpp` and `/share/vm/compiler/compileBroker.cpp`.

Most of the code is located in `/share/vm/ci/ciCacheProfiles.cpp`, a class that takes care of setting up a datastructure for the cached profiles as well as providing public methods to check if a method is cached or not. The class `/share/vm/ci/ciCacheProfilesBroker.cpp` gets called before a method that has a profile available gets compiled. It is responsible for setting up the compilation environment so the JIT compiler can use the cached profiles.

A full list of modified files and the changes can be seen in the webrev or appendix TODO.

The changes are provided in form of a patch for Hotspot version 8182 TODO. This original version is referred to as *Baseline*.

I will describe and explain the functionality and the implementation design decision in the following sections, ordered by the appearance in execution.

3.1 Creating cached profiles

The baseline version of Hotspot already offered a functionality to replay a compilation based on dumped profiling information. This is mainly used in case the JVM crashes during JIT compilation to replay the compilation again and help finding the cause of this crash. Dumping the data needed for the replay is either be done automatically in case of a crash or can be invoked manually by specifying the `DumpReplay` compile command option per method. I introduce method option called `DumpProfile` as well as a compiler flag `-XX:+DumpProfiles` that appends profiling information to a file as soon as a method gets compiled. The first option can be specified as part of the `-XX:CompileCommand` or `-XX:CompileCommandFile` flag and allows one to select single methods to

dump their profile. The second command dumps profiles of all compiled methods into a single file. The file can be opened with any text editor and is called *cached_profiles.dat*.

As soon as a method gets compiled on level 3 or level 4 all information about the methods used in the compiled method as well as their profiling information get converted to a string and written to disk. Methods compiled with level 1 and 2 will not be considered. Both are rarely used in practice and do only include none respectively little profiling information.

Since methods often get compiled multiple times and at different tiers, this can result in dumping compilation information about the same method multiple times. How this will be taken care of is described in Section 3.2. Together with some additional information about the compilation itself, for example the bytecode index of the compiled method in case of OSR, the compiler will be able to redo the same compilation on a future run of the Java virtual machine.

3.2 Initializing cached profiles

I introduce a new compiler flag `-XX:+CacheProfiles` that enables the use of profiles that have been written to disk in a previous run of the Java Virtual Machine. Per default it reads from a file called *cached_profiles.dat* but a different file can be specified using `-XX:CacheProfilesFile=other_file.dat`.

Before any cached profiles can be used the virtual machine has to parse that file and organize the profiles and compile information in a simple datastructure. This datastructure is kept in memory during the whole execution of the JVM to avoid multiple scans of the file. The parsing process gets invoked during boot up of the JVM, directly after the `CompileBroker` gets initialized. This happens before any methods get executed and blocks the JVM until finished. As mentioned in Section 3.1 the file consists of method informations, method profiles and additional compile information. The parser scans the file once and creates a so called `CompileRecord` for each of the methods that include compilation information in the file. This compile record also includes the list of method information and their profiling information. A method's compile information could have been dumped multiple times, so it can happen that there are multiple `CompileRecords` for the same method. In this case, Hotspot will only keep the `CompileRecords` that are created based on the last data written to the file. Since profiling information only grows, the compilation that happened last contains the richest profile and is considered the best. This is based on the fact that the richer the profile the more information about the method execution is known and influences the compiled version of that method. For example, a profile for a method might include data for all its branches and therefore no branches with uncommon traps which result in costly deoptimizations.

The `CompileRecord` as well as the lists of methods information and profiles are implemented as an array located in Hotspot's heap space. They get initialized with a length of 8 and grow when needed. The choice has been done for simplicity and leaves up room for further optimizations.

3.3 Using cached profiles

The idea is to use cached profiles whenever possible and if none are available continue as usual.

The thesis offers three different modes `mode0`, `mode1` and `mode2` on how the profiles are used. The following paragraph describes the behaviour of `mode0` and `mode1` and I will discuss the differences in detail in Section 3.4, especially `mode2` in Section 3.4.3.

A graphical, simplified overview of the program flow for compiling a method with the changes introduced in this thesis can be found in Figure 3.1. As mentioned before once certain thresholds

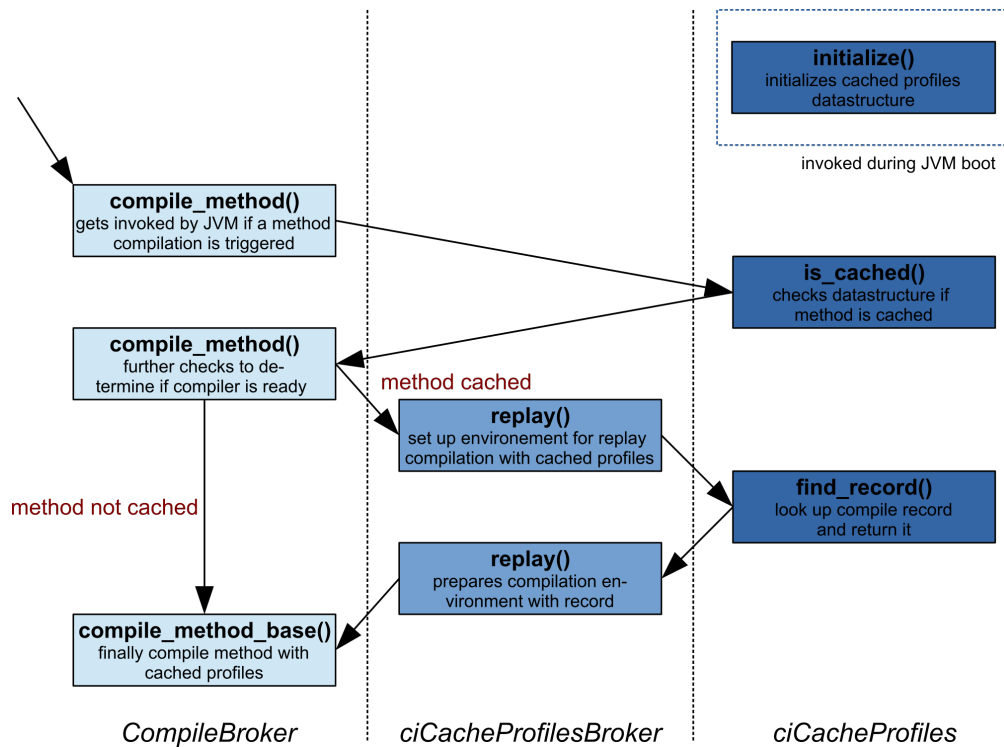


Figure 3.1: program flow for compiling a method

are exceeded a method gets scheduled for compilation. This means that the JVM will invoke a method called `compile_method()` located in the `compileBroker` class. This method for example checks if the compile queue isn't full or if there is already another compilation of that particular method running. I extended this method with a call to `ciCacheProfiles::is_cached(Method* method)` which either returns 0 if the method is not cached or returns an integer value, reflecting the compile level, in case that method has a cached profile available. Because only methods compiled with level 3 or 4 get cached, this call only gets executed if the compilation request is also of level 3 or higher. Note that this also means that if a method compilation of level 3 is initiated and a `CompileRecord` of level 4 is available that the highest level profile will be used. Therefore the method gets immediately compiled with C2 instead of C1. In case the method is not cached the execution continues like in the baseline version. Otherwise, the `compileBroker` calls into `ciCacheProfilesBroker` to replay the compilation, based on the saved profile. The

`ciCacheProfilesBroker` class then initializes the replay environment and retrieves the compile record from `ciCacheProfiles`. Subsequently the needed cached profiles get loaded to make sure they get used by a following compilation. `ciCacheProfilesBroker` then returns the execution to the `compileBroker` which continues with the steps needed to compile the method. Again some constraints are checked (e.g. if there is another compilation of the same method finished in the meantime) and a new compile job is added to the compile queue. Eventually the the method is going to be compiled using the cached profiles.

Since the implementation is basically an extension of the static class `compileBroker`, `ciCacheProfiles` and `ciCacheProfilesBroker` are static classes as well. The `compileBroker` gets invoked by the single JVM main thread and is not multi threaded, therefore there is no need to make the `compileRecord` datastructure or any of the new implementations thread safe.

3.4 Different usage modes for cached profiles

The implementation of cached profiles offers 3 different modes which differ in the transitions between the compilation tiers. The motivation as well as the advantages and disadvantages of each mode are described in the following three subsections. While `mode0` and `mode1` are similar except for the compile thresholds, `mode2` differs significantly.

3.4.1 Compile Thresholds lowered (mode 0)

The first mode is based on the consideration that a method that has a profile available does not require extensive profiling anymore. Therefore the compile thresholds (see Section 1.3) of these methods are lowered automatically. By default, they are lowered to 1% of their original values but the threshold scaling can be modified with the JVM parameter:

`-XX:CacheProfilesMode0ThresholdScaling=x.xx.`

1% results in the level 3 invocation counter being reduced from 200 to 2. This means that the method will be interpreted once but then directly trigger a compilation on the next invocation. Since the interpreter also handles class loading this decision has been made to avoid the need of doing class loading in C1 or C2 which was considered out of the scope for this thesis. As mentioned before the triggered compilation will use the latest available compile record. Eventually, most hot methods get compiled by C2 and therefore the used compiled record is usually a C2 one. In this case the JVM will jump directly from compile level 0 to compile level 4 and avoid a costly C1 compilation as well as gathering profiling information during level 0 and level 3. It will directly use the highly optimized version generated by C2 and ideally result in a lower time to reach peak performance. On the downside this increases the load on the C2 compiler and fill the compile queue more quickly. Mode 0 is the default mode and used if not further specified.

3.4.2 Unmodified Compile Thresholds (mode 1)

Mode 1 is doing exactly the same as mode 1 but does not scale the compilation thresholds automatically. This is done to decrease the load increase on C2 as mentioned in Subsection 3.4.1. Apart from this change mode 1 has the same behaviour as mode 0.

3.4.3 Modified C1 stage (mode 2)

Both modes mentioned before use cached profiles as soon as a compilation of level 3 and 4 are triggered. Since the thresholds for level 3 are smaller than the level 4 thresholds (see Appendix A.1) a method reaching a level 3 threshold could actually trigger a level 4 compilation, if the cached profile is one of level 4. So even if mode 1 is used and the thresholds are untouched C2 might get overloaded since compilations occur earlier.

Mode 2 has been designed to make as little changes as possible to the tiered compilation. and prevent C2 being more used than usual. It does so by keeping the original tiered compilation steps and compilation thresholds and compiles methods with C1 prior to C2. But since there are already profiles available there is no need to run at Tier 3 to generate full profiles but instead it uses Tier 2. Tier 2 does the same optimizations but offers only limited profiles like method invocation and backbranch counters. They are needed to know when to trigger the C2 compilation and therefore we can not use Tier 1. Tier 2 generally is considered about 30% faster than Tier 3 [5].

If then the Tier4 thresholds are reached the method is compiled using C2 and the cached profiles.

The above only makes sense if there is a C2 profile available. If only a C1 profile is available, mode 2 will only use the profile during the C1 compilation and then use the new profiles.




Figure 3.2: Tier transitions of different modes

3.5 Issues

If the profiles generated by multiple runs of the program deviate sharply it is likely that a cached profile does not fit to the current execution. In this case the compiled version would still trigger many deoptimizations and the method could end up having even worse performance since it's going to use the profile over and over again. To circumvent that behavior I modified the code that only methods which have been deoptimized less than 10 times already will get compiled using cached profiles. If they are above that limit a standard compilation will be used instead. The limit is 10 to allow a small number of recompilations. This could for example be useful when the method is deoptimized due to classes not being loaded.

3.6 Debug output

For debugging and benchmarking purposes I implemented four debug flags that can be used along with `-XX:+CacheProfiles`.

flag	description
<code>-XX:+PrintCacheProfiles</code>	enable command line debug output for cached profiles
<code>-XX:+PrintDeoptimizationCount</code>	prints amount of deoptimizations when the JVM gets shut down
<code>-XX:+PrintDeoptimizationCountVerbose</code>	prints total the amount of deoptimizations on each deoptimization
<code>-XX:+PrintCompileQueue</code>	prints the total amount of methods in the compile queue each time a method gets added

4 Performance

This section evaluates the performance of the cached profile implementation using modern benchmark suites.

4.1 Setup

To provide reliable and comparable results all tests were done on a single node of the Data Center Observatory provided by ETH [1]. A node features 2 8-Core AMD Opteron 6212 CPUs running at 2600 MHz with 128 GB of DDR3 RAM. The node is running Fedora 19 and GCC 4.8.3. All JDK builds got created on the node itself.

To compare performance the following benchmarks were used:

1. **SPECjvm 2008:** A benchmark suite developed by Standard Performance Evaluation Corporation for measuring the performance of the Java Runtime Environment [10]. I use version 2008 and I run a subset of 17 out of a total of 21 benchmarks. 4 are omitted due to incompatibility with openJDK 1.9.0.
Once finished, SPECjvm prints out the number of operations per minute. This is used to compare the performance and higher is better.
2. **Octane 2.0:** A benchmark developed by Google to measure the performance of JavaScript code found in large, real-world applications [2]. Octane runs on Nashorn, a JavaScript Engine on top of Hotspot. The version used is 2.0 and consists of 17 individual benchmarks of which 16 are used.
Octane gives each benchmark a score reflecting the performance, the higher the score, the better the performance.

The benchmarking process was automated using a number of self-written python scripts. The graphs in this chapter always show the arithmetic mean of 50 runs and the error bars display the 95% confidence intervals.

4.2 Startup performance

The main goal of cached profiles is to improve the startup performance of the JVM. Having a rich profile from an earlier execution will allow the JIT compiler to use a highly optimized version right

from the beginning. I will start by looking at SPECjvm since it offers ways to focus on the warmup. An individual description of each benchmark being used can be found in Appendix A.2.

The longer a program is running the less impact a faster warmup has. Considering most benchmarks include a warmup phase which does not count towards the final score simply running the complete benchmark suite is not an option. Instead I limited SPECjvm to 1 single operation which, depending on the benchmark take around 10 to 20 seconds. Additionally, the JVM gets restarted between each single benchmark to prevent methods shared between benchmarks being compiled already.

I run each benchmark with all cached profiling features disabled. This run is called the *baseline* and displays the current openJDK 1.9.0 performance.

I then use a single benchmark run where I dump the profiles to disk. This run is not limited to a single operation and instead uses the default values of the benchmark. By default the benchmark is limited by time and runs for about 6 minutes. The idea is that these profiles include information that are usually not available during warmup and result in less deoptimizations and better code quality.

These profiles are then used in 3 individual runs using the introduced `-XX:CacheProfiles` flag. Each run is using one of the 3 different `CacheProfilesModes`. Figures 4.1 and 4.2 shows the number




Figure 4.1: SPECjvm benchmarks on all different modes

of operations per minute, measured for each benchmark individually. The




Figure 4.2: SPECjvm scimark benchmarks on all different modes

4.3 Deoptimizations

4.4 Effect on compile queue




Figure 4.3:




Figure 4.4:




Figure 4.5:




Figure 4.6:

5 Possible Improvements

- Better datastructure - Merging multiple profiles cleverly

6 Conclusion

A Appendix

A.1 Tiered Compilation Thresholds

flag	description	default
CompileThresholdScaling	number of interpreted method invocations before (re-)compiling	1.0
Tier0InvokeNotifyFreqLog	Interpreter (tier 0) invocation notification frequency	7
Tier2InvokeNotifyFreqLog	C1 without MDO (tier 2) invocation notification frequency	11
Tier3InvokeNotifyFreqLog	C1 with MDO profiling (tier 3) invocation notification frequency	10
Tier23InlineeNotifyFreqLog	Inlinee invocation (tiers 2 and 3) notification frequency	20
Tier0BackedgeNotifyFreqLog	Interpreter (tier 0) invocation notification frequency	10
Tier2BackedgeNotifyFreqLog	C1 without MDO (tier 2) invocation notification frequency	14
Tier3BackedgeNotifyFreqLog	C1 with MDO profiling (tier 3) invocation notification frequency	13
Tier2CompileThreshold	threshold at which tier 2 compilation is invoked	0
Tier2BackEdgeThreshold	Back edge threshold at which tier 2 compilation is invoked	0
Tier3InvocationThreshold	Compile if number of method invocations crosses this threshold	200
Tier3MinInvocationThreshold	Minimum invocation to compile at tier 3	100
Tier3CompileThreshold	Threshold at which tier 3 compilation is invoked (invocation minimum must be satisfied)	2000
Tier3BackEdgeThreshold	Back edge threshold at which tier 3 OSR compilation is invoked	60000
Tier4InvocationThreshold	Compile if number of method invocations crosses this threshold	5000
Tier4MinInvocationThreshold	Minimum invocation to compile at tier 4	600
Tier4CompileThreshold	Threshold at which tier 4 compilation is invoked (invocation minimum must be satisfied)	15000
Tier4BackEdgeThreshold	Back edge threshold at which tier 4 OSR compilation is invoked	40000

A.2 SPECjvm Benchmark

This list gives a short description of the benchmarks that are part of the SPECjvm 2008 Benchmark Suite. The list is directly taken from <https://www.spec.org/jvm2008/docs/benchmarks/index.html> and put in as a reference.

- **Compress:** This benchmark compresses data, using a modified Lempel-Ziv method (LZW). Basically finds common substrings and replaces them with a variable size code. This is deterministic, and can be done on the fly. Thus, the decompression procedure needs no input table, but tracks the way the table was built. Algorithm from "A Technique for High

Performance Data Compression”, Terry A. Welch, IEEE Computer Vol 17, No 6 (June 1984), pp 8-19.

This is a Java port of the 129.compress benchmark from CPU95, but improves upon that benchmark in that it compresses real data from files instead of synthetically generated data as in 129.compress.

- **Crypto:** This benchmark focuses on different areas of crypto and are split in three different sub-benchmarks. The different benchmarks use the implementation inside the product and will therefore focus on both the vendor implementation of the protocol as well as how it is executed.
 - **aes** encrypt and decrypt using the AES and DES protocols, using CBC/PKCS5Padding and CBC/NoPadding. Input data size is 100 bytes and 713 kB.
 - **rsa** encrypt and decrypt using the RSA protocol, using input data of size 100 bytes and 16 kB.
 - **signverify** sign and verify using MD5withRSA, SHA1withRSA, SHA1withDSA and SHA256withRSA protocols. Input data size of 1 kB, 65 kB and 1 MB.
- **Derby:** This benchmark uses an open-source database written in pure Java. It is synthesized with business logic to stress the BigDecimal library. It is a direct replacement to the SPECjvm98 db benchmark but is more capable and represents as close to a "real" application. The focus of this benchmark is on BigDecimal computations (based on telco benchmark) and database logic, especially, on locks behavior. BigDecimal computations are trying to be outside 64-bit to examine not only 'simple' BigDecimal, where 'long' is used often for internal representation.
- **MPEGaudio:** This benchmark is very similar to the SPECjvm98 mpegaudio. The mp3 library has been replaced with JLayer, an LGPL mp3 library. Its floating-point heavy and a good test of mp3 decoding. Input data were taken from SPECjvm98.
- **Scimark:** This benchmark was developed by NIST and is widely used by the industry as a floating point benchmark. Each of the subtests (**fft**, **lu**, **monte_carlo**, **sor**, **sparse**) were incorporated into SPECjvm2008. There are two versions of this test, one with a `largeDataset` (32Mbytes) which stresses the memory subsystem and a `smallDataset` which stresses the JVMs (512Kbytes).
- **Serial:** This benchmark serializes and deserializes primitives and objects, using data from the JBoss benchmark. The benchmark has a producer-consumer scenario where serialized objects are sent via sockets and deserialized by a consumer on the same system. The benchmark heavily stress the `Object.equals()` test.
- **Sunflow:** This benchmark tests graphics visualization using an open source, internally multi-threaded global illumination rendering system. The sunflow library is threaded internally, i.e. it's possible to run several bundles of dependent threads to render an image. The

number of internal sunflow threads is required to be 4 for a compliant run. It is however possible to configure in property `specjvm.benchmark.sunflow.threads.per.instance`, but no more than 16, per sunflow design. Per default, the benchmark harness will use half the number of benchmark threads, i.e. will run as many sunflow benchmark instances in parallel as half the number of hardware threads. This can be configured in `specjvm.benchmark.threads.sunflow`.

- **XML:** This benchmark has two sub-benchmarks: `XML.transform` and `XML.validation`. `XML.transform` exercises the JRE's implementation of `javax.xml.transform` (and associated APIs) by applying style sheets (.xsl files) to XML documents. The style sheets and XML documents are several real life examples that vary in size (3KB to 156KB) and in the style sheet features that are used most heavily. One "operation" of `XML.transform` consists of processing each style sheet / document pair, accessing the XML document as a DOM source, a SAX source, and a Stream source. In order that each style sheet / document pair contribute about equally to the time taken for a single operation, some of the input pairs are processed multiple times during one operation.

Result verification for `XML.transform` is somewhat more complex than for other of the benchmarks because different XML style sheet processors can produce results that are slightly different from each other, but all still correct. In brief, the process used is this. First, before the measurement interval begins the workload is run once and the output is collected, canonicalized (per the specification of canonical XML form) and compared with the expected canonicalized output. Output from transforms that produce HTML is converted to XML before canonicalization. Also, a checksum is generated from this output. Inside the measurement interval the output from each operation is only checked using the checksum.

`XML.validation` exercises the JRE's implementation of `javax.xml.validation` (and associated APIs) by validating XML instance documents against XML schemata (.xsd files). The schemata and XML documents are several real life examples that vary in size (1KB to 607KB) and in the XML schema features that are used most heavily. One "operation" of `XML.validation` consists of processing each style sheet / document pair, accessing the XML document as a DOM source and a SAX source. As in `XML.transform`, some of the input pairs are processed multiple times during one operation so that each input pair contributes about equally to the time taken for a single operation.

A.3 Octane Benchmark

What follows is an overview of the benchmarks Octane consists of. The original list can be found on <https://developers.google.com/octane/benchmark>.

- **Richards:** OS kernel simulation benchmark, originally written in BCPL by Martin Richards (539 lines).
 - Main focus: *property load/store, function/method calls*
 - Secondary focus: *code optimization, elimination of redundant code*

- **Deltablue:** One-way constraint solver, originally written in Smalltalk by John Maloney and Mario Wolczko (880 lines).
 - Main focus: *polymorphism*
 - Secondary focus: *OO-style programming*
- **Raytrace:** Ray tracer benchmark based on code by Adam Burmister (904 lines).
 - Main focus: *argument object, apply*
 - Secondary focus: *prototype library object, creation pattern*
- **Regexp:** Regular expression benchmark generated by extracting regular expression operations from 50 of the most popular web pages (1761 lines).
 - Main focus: *Regular expressions*
- **NavierStokes:** 2D NavierStokes equations solver, heavily manipulates double precision arrays. Based on Oliver Hunt's code (387 lines).
 - Main focus: *reading and writing numeric arrays.*
 - Secondary focus: *floating point math.*
- **Crypto:** Encryption and decryption benchmark based on code by Tom Wu (1698 lines).
 - Main focus: *bit operations*
- **Splay:** Data manipulation benchmark that deals with splay trees and exercises the automatic memory management subsystem (394 lines).
 - Main focus: *Fast object creation, destruction*
- **SplayLatency:** The Splay test stresses the Garbage Collection subsystem of a VM. Splay-Latency instruments the existing Splay code with frequent measurement checkpoints. A long pause between checkpoints is an indication of high latency in the GC. This test measures the frequency of latency pauses, classifies them into buckets and penalizes frequent long pauses with a low score.
 - Main focus: *Garbage Collection latency*
- **EarleyBoyer:** Classic Scheme benchmarks, translated to JavaScript by Florian Loitsch's Scheme2Js compiler (4684 lines).
 - Main focus: *Fast object creation, destruction*
 - Secondary focus: *closures, arguments object*
- **pdf.js:** Mozilla's PDF Reader implemented in JavaScript. It measures decoding and interpretation time (33,056 lines).

- Main focus: *array and typed arrays manipulations*.
- Secondary focus: *math and bit operations, support for future language features (e.g. promises)*
- **Mandree1:** Runs the 3D Bullet Physics Engine ported from C++ to JavaScript via Mandree1 (277,377 lines).
 - Main focus: *emulation*
- **Mandree1Latency:** Similar to the SplayLatency test, this test instruments the Mandree1 benchmark with frequent time measurement checkpoints. Since Mandree1 stresses the VM's compiler, this test provides an indication of the latency introduced by the compiler. Long pauses between measurement checkpoints lower the final score.
 - Main focus: *Compiler latency*
- **GB Emulator:** Emulates the portable console's architecture and runs a demanding 3D simulation, all in JavaScript (11,097 lines).
 - Main focus: *emulation*
- **Code loading:** Measures how quickly a JavaScript engine can start executing code after loading a large JavaScript program, social widget being a common example. The source for this test is derived from open source libraries (Closure, jQuery) (1,530 lines).
 - Main focus: *JavaScript parsing and compilation*
- **Box2DWeb:** Based on Box2DWeb, the popular 2D physics engine originally written by Erin Catto, ported to JavaScript. (560 lines, 9000+ de-minified)
 - Main focus: *floating point math*.
 - Secondary focus: *properties containing doubles, accessor properties*.
- **Typescript:** Microsoft's TypeScript compiler is a complex application. This test measures the time TypeScript takes to compile itself and is a proxy of how well a VM handles complex and sizable Javascript applications (25,918 lines).
 - Main focus: *run complex, heavy applications*

Bibliography

- [1] ETH. ETH Data Center Observatory. <https://wiki.dco.ethz.ch/bin/viewauth/Main/Cluster>, 2015.
- [2] Google. Octane 2.0 Benchmark. <https://developers.google.com/octane/>, 2015.
- [3] T. Hartmann, A. Noll, and T. R. Gross. Efficient code management for dynamic multi-tiered compilation systems. In *2014 International Conference on Principles and Practices of Programming on the Java Platform Virtual Machines, Languages and Tools, PPPJ '14, Cracow, Poland, September 23-26, 2014*, pages 51–62, 2014.
- [4] V. Ivanov. JIT-compiler in JVM seen by a Java developer. <http://www.stanford.edu/class/cs343/resources/java-hotspot.pdf>, 2013.
- [5] Oracle Corporation. Code for advancedThresholdPolicy.hpp. <http://hg.openjdk.java.net/jdk9/hs-comp/hotspot/file/63337cc98898/src/share/vm/runtime/advancedThresholdPolicy.hpp>, 2013.
- [6] Oracle Corporation. hprof - A heap/CPU profiling tool. <https://docs.oracle.com/javase/7/docs/technotes/samples/hprof.html>, 2014.
- [7] Oracle Corporation. javac - Java programming language compiler. <https://docs.oracle.com/javase/7/docs/technotes/tools/windows/javac.html>, 2014.
- [8] M. Paleczny, C. Vick, and C. Click. The Java HotSpot™Server Compiler. https://www.usenix.org/legacy/events/jvm01/full_papers/paleczny/paleczny.pdf, 2001. Paper from JVM '01.
- [9] T. Rodriguez and K. Russell. Client Compiler for the Java HotSpot™Virtual Machine: Technology and Application. <http://www.oracle.com/technetwork/java/javase/tech/3198-d1-150056.pdf>, 2002. Talk from JavaOne 2002.
- [10] Standard Performance Evaluation Corporation. SPECjvm2008 Benchmark. <https://www.spec.org/jvm2008/>, 2008.